

Huseyin Can Minareci - Noam Shmuel

June 10, 2020

0.1 Faculty of Economic Sciences University of Warsaw

Machine Learning 1: Classification Methods

Dr. Maciej Wilamowski

1 Used Car Price Prediction Analysis

Huseyin Can Minareci & Noam Shmuel

2 Introduction

In this project, we aim to find the best regression model for dataset in order to be able to predict used cars prices.

We have used public dataset from Kaggle which was scraped with Scrapy by Orges Leka from Ebay-Kleinanzeigen. It contains more than 370000 cars and 20 features of each of them like price, power, model etc...

Original dataset had plenty missing values, outliers and mistakes, to be able to work on it we had to clean, impute and change their format.

We select the features carefully and extract new features from raw data.

In order to achieve this goal we used Linear Model, K Nearest Neighbor and Random Forest.

3 Data Preparation

```
[1]: # Importing the packages which we are going to use
import pandas as pd
import numpy as np
import math
import random
import time
import matplotlib
from math import sqrt
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import datasets, linear_model, preprocessing, svm, neighbors,
↪ metrics
from sklearn.preprocessing import StandardScaler, Normalizer, MinMaxScaler
from datetime import datetime
```

```

import statsmodels.api as sm
import statsmodels.formula.api as smf
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import KFold, train_test_split, GridSearchCV

%matplotlib inline

```

```

[2]: # reading dataset
df = pd.read_csv('autos.csv', sep=',', header=0, encoding='cp1252')
df.sample(5)

```

```

[2]:
      dateCrawled      name \
300022  2016-03-25 18:46:09 Opel_Astra_2.0_DTI_Caravan_Elegance
326791  2016-03-17 20:41:51 Ford_Fiesta_1.4_Ghia
220614  2016-03-20 14:47:23 BMW_E46_318i_Touring_143PS_PDC_Klimaautomatik
224538  2016-03-23 09:55:40 Ford_Focus_1.4_16V
157406  2016-03-31 15:53:13 Volkswagen_Touran_1.9_TDI_Comfortline

      seller offerType  price  abtest vehicleType  yearOfRegistration \
300022  privat  Angebot   1690   test      kombi                2003
326791  privat  Angebot   5500 control  kleinwagen                2009
220614  privat  Angebot   3800   test      kombi                2004
224538  privat  Angebot   4499   test  limousine                2007
157406  privat  Angebot   4899 control      bus                2004

      gearbox  powerPS  model  kilometer  monthOfRegistration  fuelType \
300022  manuell     101  astra    150000                6  diesel
326791  manuell     97  fiesta   125000                3  benzin
220614  manuell    143   3er    150000                1  benzin
224538  manuell     80  focus   100000                8  benzin
157406  manuell    105  touran   150000                8  diesel

      brand notRepairedDamage  dateCreated  nrOfPictures \
300022   opel              nein  2016-03-25 00:00:00        0
326791   ford              nein  2016-03-17 00:00:00        0
220614   bmw              nein  2016-03-20 00:00:00        0
224538   ford              nein  2016-03-23 00:00:00        0
157406 volkswagen          nein  2016-03-31 00:00:00        0

      postalCode  lastSeen
300022     90602  2016-04-07 00:17:10
326791     41468  2016-04-07 09:15:50
220614     76131  2016-04-06 17:47:01
224538     22946  2016-04-07 10:16:30

```

157406 67227 2016-04-06 08:46:23

The variables explained as follows: price : the price on the advert to sell the car, This is the dependent variable in all of the upcoming models.

dateCrawled : when advert was first crawled, all field-values are taken from this date

name : headline, which the owner of the car gave to the advert

seller : who is selling the car(private or dealer)

offerType : offer(car to sell) or request(car to buy)

abtest : ebay-intern variable (a/b testing)

vehicleType : one of eight vehicle-categories

yearOfRegistration : at which year the car was first registered

gearbox : manual or automatic

powerPS : the power of the car in PS

model : the car's model

kilometer : how many kilometres the car has driven

monthOfRegistration : at which month the car was first registered

fuelType : one of seven fuel-categories

brand : the car's brand

notRepairedDamage : if the car has a damage which is not repaired yet

dateCreated : the date for which the advert at 'ebay Kleinanzeigen' was created

nrOfPictures : number of pictures in the advert

postalCode : where in germany the car is located

lastSeenOnline : when the crawler saw this advert last online

3.0.1 Duplicates and outliers

Since data is crawled by scraper we may have some duplicates it will be better to check and drop them as first step.

```
[3]: # Removing the duplicates
df1 = df.drop_duplicates(['name', 'price', 'vehicleType', 'yearOfRegistration'
                           ↵
                           ↪, 'gearbox', 'powerPS', 'model', 'kilometer', 'monthOfRegistration', 'fuelType'
                           , 'notRepairedDamage'])
left = 100 * df1['name'].count() / df['name'].count()
print("The amount of data left:", left, "%")
```

The amount of data left: 97.64270795202515 %

```
[4]: df.describe()
```

```
[4]:           price  yearOfRegistration  powerPS  kilometer \
count  3.715280e+05      371528.000000  371528.000000  371528.000000
```

mean	1.729514e+04	2004.577997	115.549477	125618.688228
std	3.587954e+06	92.866598	192.139578	40112.337051
min	0.000000e+00	1000.000000	0.000000	5000.000000
25%	1.150000e+03	1999.000000	70.000000	125000.000000
50%	2.950000e+03	2003.000000	105.000000	150000.000000
75%	7.200000e+03	2008.000000	150.000000	150000.000000
max	2.147484e+09	9999.000000	20000.000000	150000.000000

	monthOfRegistration	nrOfPictures	postalCode
count	371528.000000	371528.0	371528.000000
mean	5.734445	0.0	50820.66764
std	3.712412	0.0	25799.08247
min	0.000000	0.0	1067.000000
25%	3.000000	0.0	30459.000000
50%	6.000000	0.0	49610.000000
75%	9.000000	0.0	71546.000000
max	12.000000	0.0	99998.000000

It can be clearly seen from summary statistics that there are outliers(yearOfRegistration, powerPS, price) in the data. We should remove them in order to continue.

```
[5]: print("Cars more expensive than 40000:", (dfl.price > 40000).sum())
      print("Cars cheaper than 250:", (dfl.price < 250).sum())
      print("Cars which has less than 10PS:" , (dfl.powerPS < 10).sum())
      print("Cars which has more than 400PS:" , (dfl.powerPS > 400).sum())
      print("Cars newer than 2017:" , (dfl.yearOfRegistration >= 2017).sum())
      print("Cars older than 1990:" , (dfl.yearOfRegistration < 1990).sum())
```

```
Cars more expensive than 40000: 2692
Cars cheaper than 250: 19530
Cars which has less than 10PS: 39579
Cars which has more than 400PS: 1864
Cars newer than 2017: 14483
Cars older than 1990: 10495
```

There are only 2692 cars which are more expensive than 40000 since we have over 370000 cars we can easily call them as outliers and remove them from our dataset. Also we don't want to take into account the obvious mistakes like cars cheaper than 250.

Looks like we have a lot of cars which has 0 PS since something like that it is not possible we need to remove them from our dataset as well.

```
[6]: # Removing the outliers
      dfl = dfl[
          (dfl.yearOfRegistration <= 2017)
          & (dfl.yearOfRegistration >= 1990)
          & (dfl.price >= 250)
          & (dfl.powerPS > 0)]
```

```

    & (df1.powerPS >= 10)
    & (df1.powerPS <= 400)]
left = 100 * df1['name'].count() / df['name'].count()
print("The amount of data left:", left, "%")

```

The amount of data left: 79.97378394091427 %

We check again summary statistics to see if outliers removed from dataset succesfully.

```
[7]: df1.describe()
```

```

[7]:          price  yearOfRegistration  powerPS  kilometer \
count  297125.000000    297125.000000  297125.000000  297125.000000
mean    5805.787638      2004.062714    124.808094  126377.063525
std     6364.402007        5.932930    55.794467   38409.786388
min      250.000000      1990.000000    10.000000   5000.000000
25%     1450.000000      2000.000000    82.000000  125000.000000
50%     3499.000000      2004.000000   116.000000  150000.000000
75%     7800.000000      2008.000000   150.000000  150000.000000
max     40000.000000     2017.000000   400.000000  150000.000000

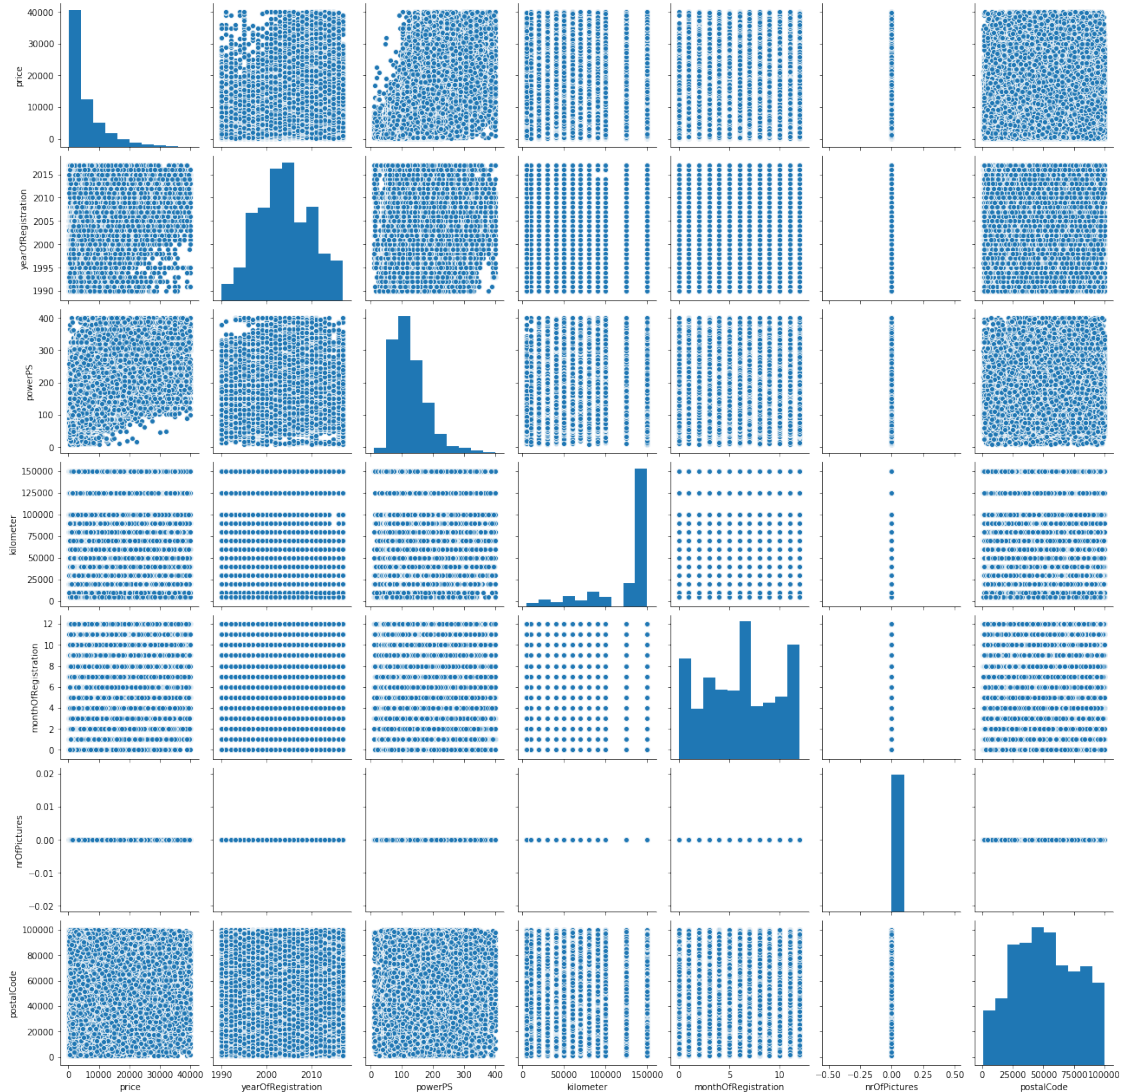
          monthOfRegistration  nrOfPictures  postalCode
count      297125.000000      297125.0  297125.000000
mean         6.016505          0.0    51641.176182
std          3.577724          0.0    25686.994168
min           0.000000          0.0     1067.000000
25%           3.000000          0.0    31246.000000
50%           6.000000          0.0    50765.000000
75%           9.000000          0.0    72393.000000
max          12.000000          0.0    99998.000000

```

3.0.2 Exploratory Data Analysis

```
[8]: sns.pairplot(df1)
```

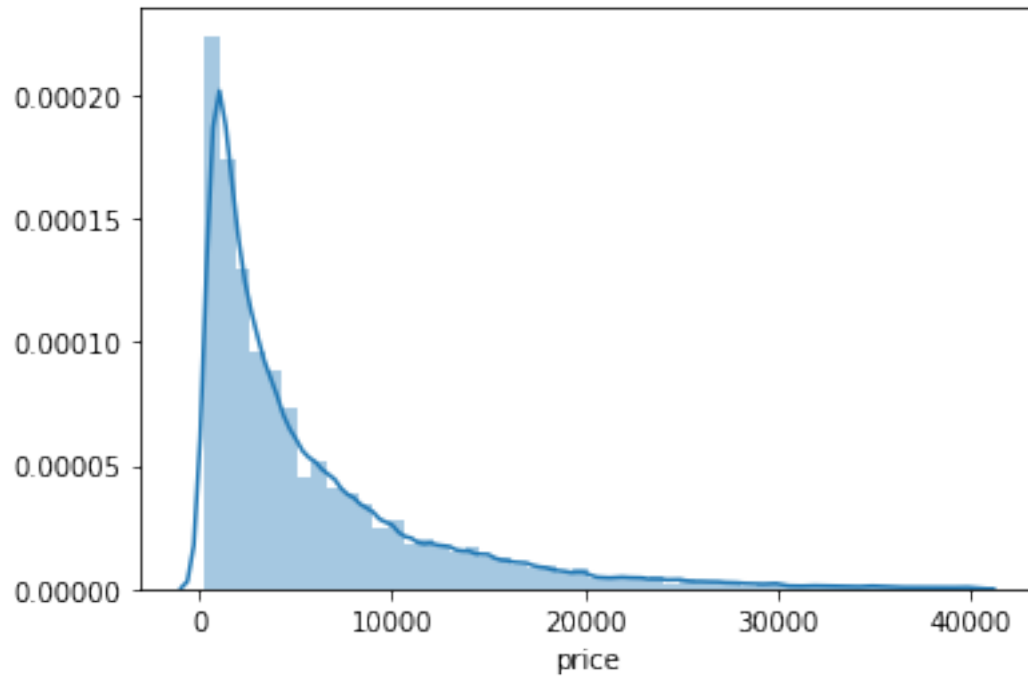
```
[8]: <seaborn.axisgrid.PairGrid at 0x1ba0d39cc50>
```



From pairwise plots we can see how our variables looks like and go deeper to the necessary ones.

```
[9]: sns.distplot(df1["price"])
```

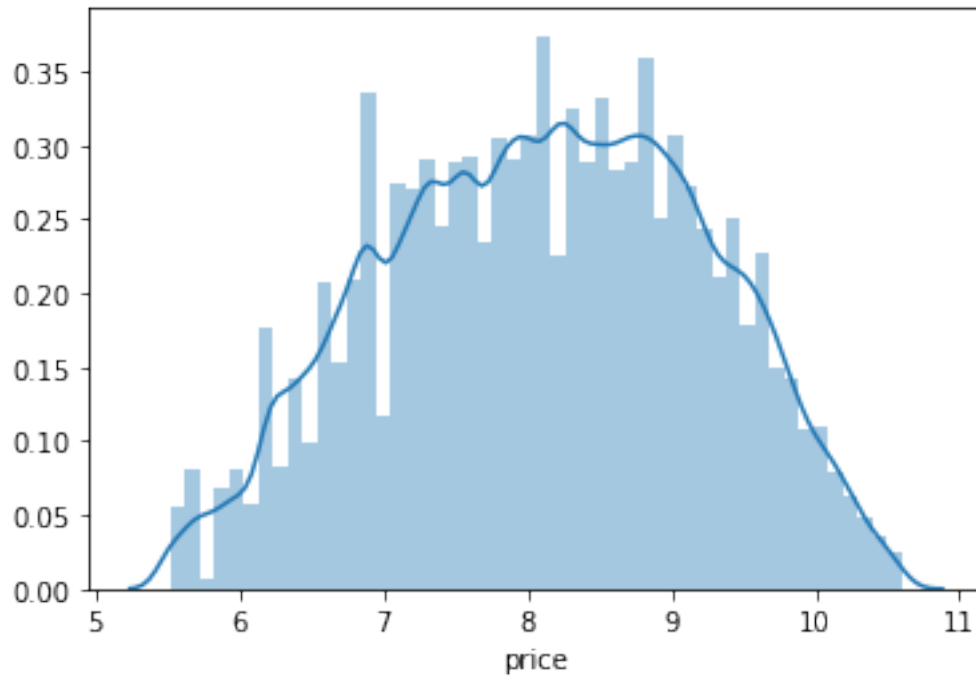
```
[9]: <matplotlib.axes._subplots.AxesSubplot at 0x1ba185b3128>
```



We have make a log transformation on price in in order to get rid of right skewness.

```
[10]: dfl['priceNormal'] = dfl['price']  
      dfl['price'] = np.log(dfl['price'])  
      sns.distplot(dfl["price"])
```

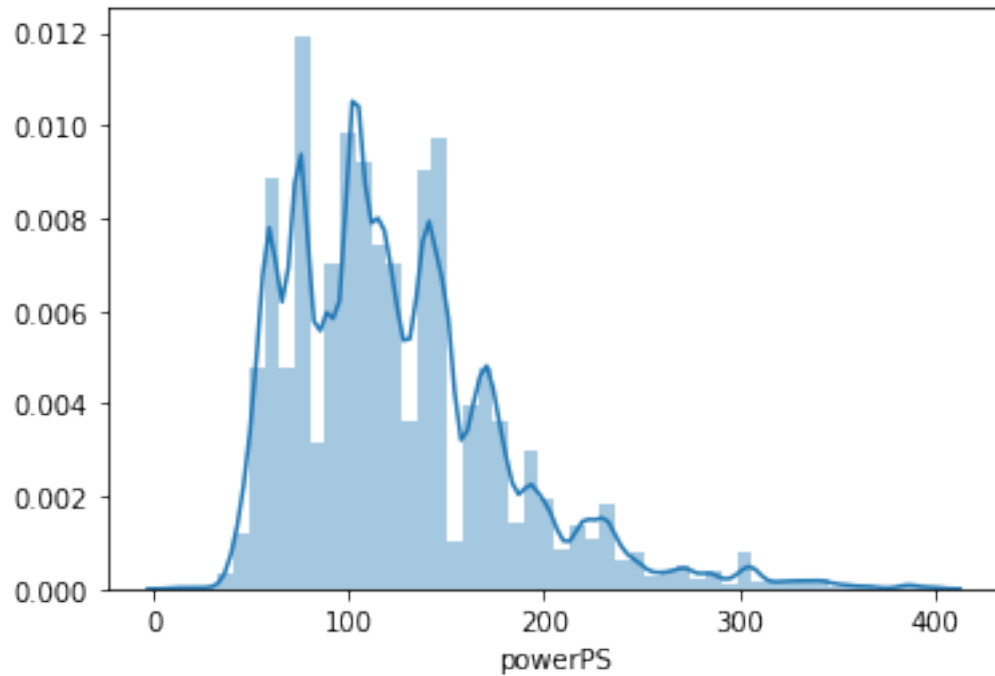
```
[10]: <matplotlib.axes._subplots.AxesSubplot at 0x1ba186be4e0>
```



Since we did log transformation to our dependent variable we need to do also to the skewed independent variables

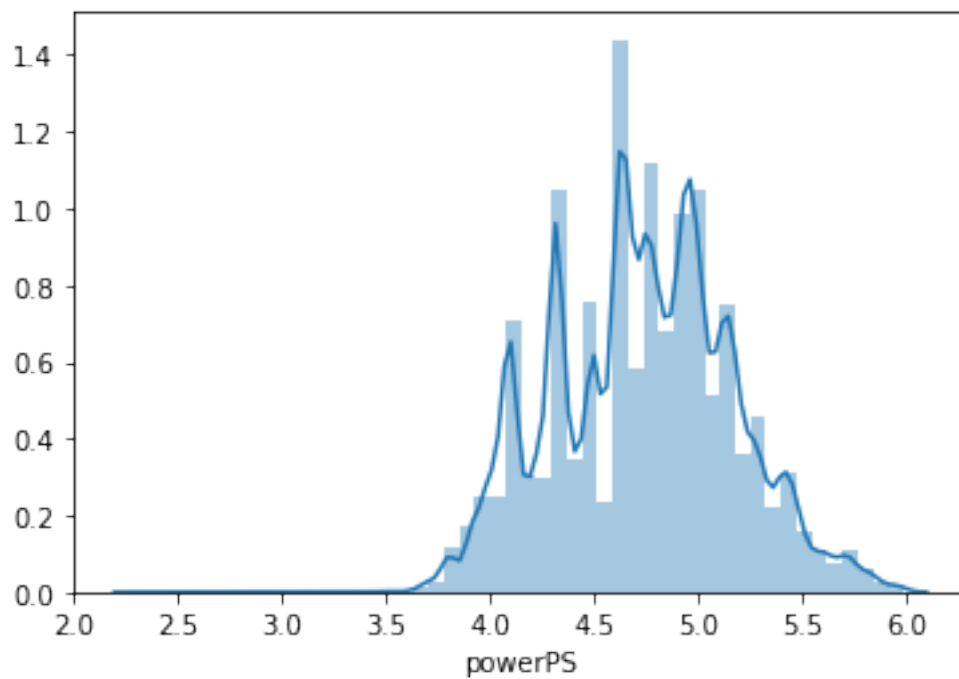
```
[11]: sns.distplot(df1["powerPS"])
```

```
[11]: <matplotlib.axes._subplots.AxesSubplot at 0x1ba187a5390>
```

```
[12]: df1['powerPS'] = np.log(df1['powerPS'])
      sns.distplot(df1["powerPS"])
```

```
[12]: <matplotlib.axes._subplots.AxesSubplot at 0x1ba18885320>
```



3.1 Feature Engineering & Selection

After dealing with duplicates, outliers and transformation, we moved to the next part; Feature Engineering.

3.1.1 daysBeforeSold

```
[13]: df1.loc[:, ("dateCrawled", "dateCreated", "postalCode", "lastSeen")]
```

```
[13]:
```

	dateCrawled	dateCreated	postalCode	\
1	2016-03-24 10:58:45	2016-03-24 00:00:00	66954	
2	2016-03-14 12:52:21	2016-03-14 00:00:00	90480	
3	2016-03-17 16:54:04	2016-03-17 00:00:00	91074	
4	2016-03-31 17:25:20	2016-03-31 00:00:00	60437	
5	2016-04-04 17:36:23	2016-04-04 00:00:00	33775	
...	
371520	2016-03-19 19:53:49	2016-03-19 00:00:00	96465	
371524	2016-03-05 19:56:21	2016-03-05 00:00:00	26135	
371525	2016-03-19 18:57:12	2016-03-19 00:00:00	87439	
371526	2016-03-20 19:41:08	2016-03-20 00:00:00	40764	
371527	2016-03-07 19:39:19	2016-03-07 00:00:00	73326	
	lastSeen			
1	2016-04-07 01:46:50			
2	2016-04-05 12:47:46			
3	2016-03-17 17:40:17			
4	2016-04-06 10:17:21			
5	2016-04-06 19:17:07			
...	...			
371520	2016-03-19 20:44:43			
371524	2016-03-11 18:17:12			
371525	2016-04-07 07:15:26			
371526	2016-03-24 12:45:21			
371527	2016-03-22 03:17:10			

```
[297125 rows x 4 columns]
```

We don't think "dateCrawled" and "postalCode" would be useful for price prediction so we decided to remove them.

Although "dateCreated" and "lastSeen" could be useful if can be modified as how many days advert was on the site before it sold.(or removed)

```
[14]: df1 = df1.drop(["dateCrawled", "postalCode"], axis='columns')
```

In the dataset dateCreated and lastSeen is in string format. In order get the days we choosed to

change the format to timestamp.

Since we can use subtraction in timestamp, we subtract dateCreated from lastSeen.

And finally take just take the day since we don't need hours minutes and seconds in our models.

```
[15]: # creating funtion to change string to timestamp
def changeformat(date):
    datefr = datetime.strptime(date, '%Y-%m-%d %H:%M:%S')
    return datefr
```

```
[16]: df1["dateCreated"] = df1["dateCreated"].apply(changeformat)
df1["lastSeen"] = df1["lastSeen"].apply(changeformat)
df1["diff"] = df1["lastSeen"]-df1["dateCreated"]
df1["daysBeforeSold"] = df1["diff"].apply((lambda x: x.days))
```

```
[17]: df1.loc[:,("dateCreated", "lastSeen", "diff", "daysBeforeSold")]
```

```
[17]:
```

	dateCreated	lastSeen	diff	daysBeforeSold
1	2016-03-24	2016-04-07 01:46:50	14 days	14
2	2016-03-14	2016-04-05 12:47:46	22 days	22
3	2016-03-17	2016-03-17 17:40:17	0 days	0
4	2016-03-31	2016-04-06 10:17:21	6 days	6
5	2016-04-04	2016-04-06 19:17:07	2 days	2
...
371520	2016-03-19	2016-03-19 20:44:43	0 days	0
371524	2016-03-05	2016-03-11 18:17:12	6 days	6
371525	2016-03-19	2016-04-07 07:15:26	19 days	19
371526	2016-03-20	2016-03-24 12:45:21	4 days	4
371527	2016-03-07	2016-03-22 03:17:10	15 days	15

[297125 rows x 4 columns]

After successfully getting the day difference in integer format we removed “dateCreated”, “lastSeen” and “diff” from the data.

```
[18]: df1 = df1.drop(["dateCreated", "lastSeen", "diff"], axis='columns')
```

3.1.2 namelen

```
[19]: df1["name"]
```

```
[19]:
```

1	A5_Sportback_2.7_Tdi
2	Jeep_Grand_Cherokee_"Overland"
3	GOLF_4_1_4__3TÜRER
4	Skoda_Fabia_1.4_TDI_PD_Classic
5	BMW_316i___e36_Limousine__Bastlerfahrzeug_Ex...
...	...
371520	turbo_defekt

```

371524          Smart_smart_leistungssteigerung_100ps
371525          Volkswagen_Multivan_T4_TDI_7DC_UY2
371526          VW_Golf_Kombi_1_91_TDI
371527          BMW_M135i_vollausgestattet_NP_52.720___Euro
Name: name, Length: 297125, dtype: object

```

Names of adverts are not usefull how they are right now we want to try to get something from this.

Taking their length and checking if it would be helpful in our prediction sounded like it worths to give it a try.

```
[20]: df1['namelen'] = [min(70, len(n)) for n in df1['name']]
```

```
[21]: df1.loc[:, 'namelen']
```

```

[21]: 1          20
      2          30
      3          18
      4          30
      5          50
      ..
371520      12
371524      37
371525      34
371526      22
371527      44
Name: namelen, Length: 297125, dtype: int64

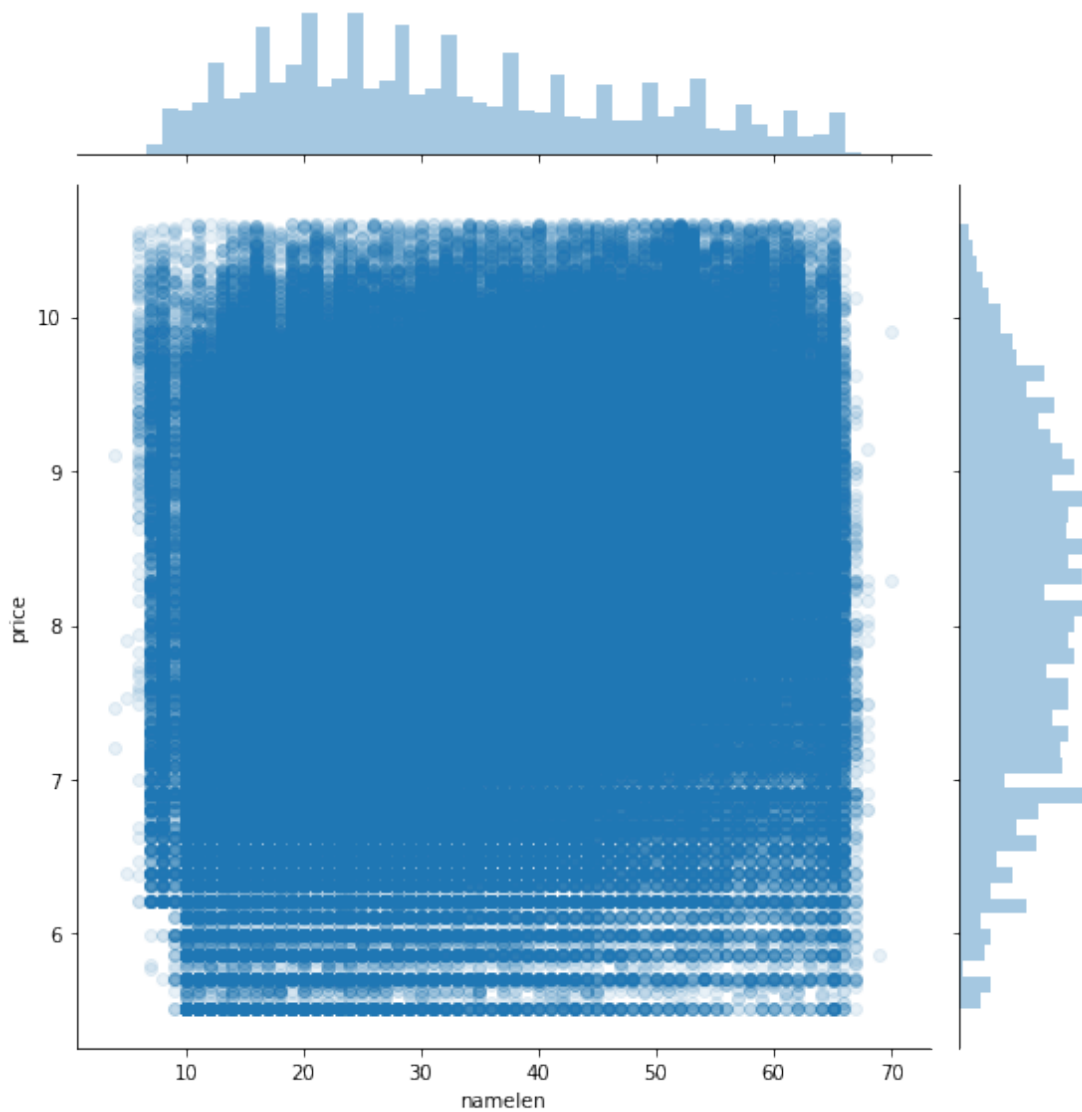
```

```

[22]: sns.jointplot(x='namelen',
                    y='price',
                    data=df1[['namelen', 'price']],
                    alpha=0.1,
                    height=8)

```

```
[22]: <seaborn.axisgrid.JointGrid at 0x1ba1ae40b38>
```

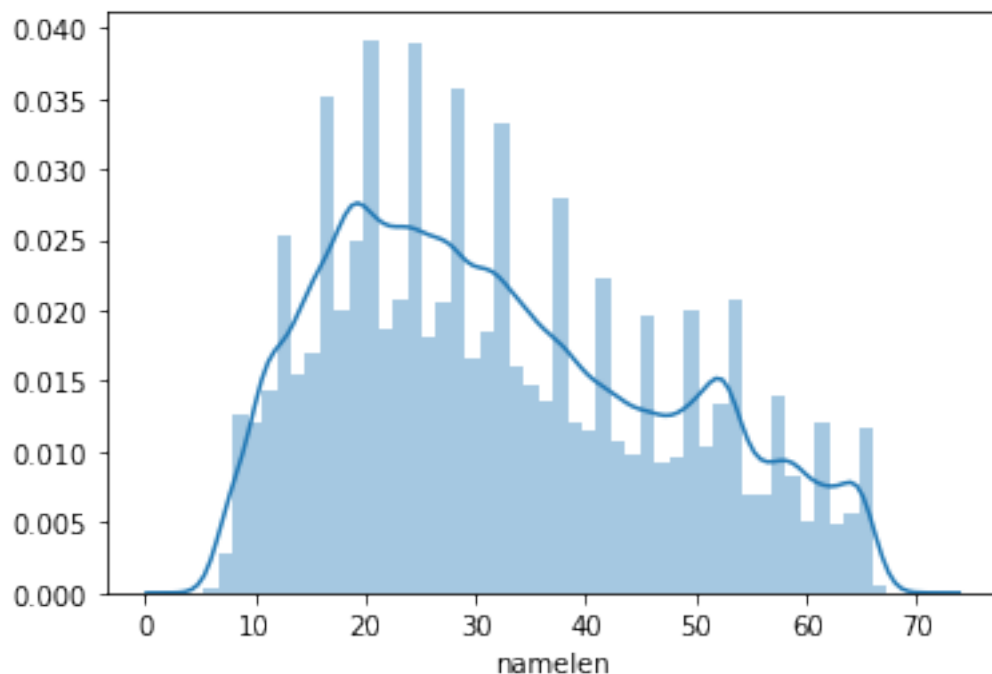


It seems that a name length can help us to predict the price since it looks like there is some connection between them.

Longer the name is more expensive the car, it can be because the cars with additional features would take longer to describe them. From other side shorter the explanation can lead us to think that car doesn't have much to write about that's why it's cheaper.

```
[23]: sns.distplot(df1.namelen)
```

```
[23]: <matplotlib.axes._subplots.AxesSubplot at 0x1ba1ae406d8>
```



We dropped the name column since we took length of it.

```
[24]: df1 = df1.drop(["name"], axis='columns')
```

3.1.3 Age

Instead of dealing with years we create “age” of the car by using yearOfRegistration and simply subtracting it from 2017(the year data collected)

```
[25]: df1["age"] = df1["yearOfRegistration"].apply((lambda x: max(0, 2017-x)))
```

```
[26]: df1.describe()
```

```
[26]:
```

	price	yearOfRegistration	powerPS	kilometer \
count	297125.000000	297125.000000	297125.000000	297125.000000
mean	8.107775	2004.062714	4.733331	126377.063525
std	1.110681	5.932930	0.433482	38409.786388
min	5.521461	1990.000000	2.302585	5000.000000
25%	7.279319	2000.000000	4.406719	125000.000000
50%	8.160232	2004.000000	4.753590	150000.000000
75%	8.961879	2008.000000	5.010635	150000.000000
max	10.596635	2017.000000	5.991465	150000.000000

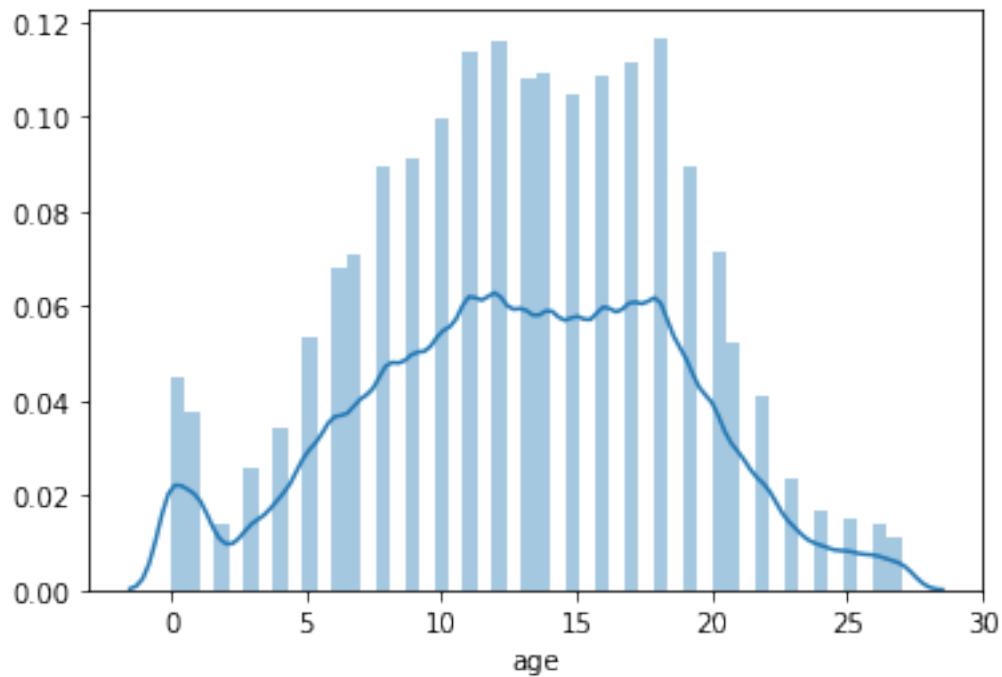
	monthOfRegistration	nrOfPictures	priceNormal	daysBeforeSold \
count	297125.000000	297125.0	297125.000000	297125.000000

mean	6.016505	0.0	5805.787638	9.059598
std	3.577724	0.0	6364.402007	8.629988
min	0.000000	0.0	250.000000	0.000000
25%	3.000000	0.0	1450.000000	2.000000
50%	6.000000	0.0	3499.000000	6.000000
75%	9.000000	0.0	7800.000000	14.000000
max	12.000000	0.0	40000.000000	384.000000

	namelen	age
count	297125.000000	297125.000000
mean	32.280074	12.937286
std	15.048183	5.932930
min	4.000000	0.000000
25%	20.000000	9.000000
50%	30.000000	13.000000
75%	43.000000	17.000000
max	70.000000	27.000000

```
[27]: sns.distplot(df1["age"])
```

```
[27]: <matplotlib.axes._subplots.AxesSubplot at 0x1ba23b14da0>
```



Since we have age of the car we drop yearOfRegistration and monthOfRegistration.

```
[28]: df1 = df1.drop(["yearOfRegistration", "monthOfRegistration"], axis='columns')
```

3.1.4 gearbox abtest notRepairedDamage

```
[29]: print(df1.gearbox.value_counts())  
      print(df1.abtest.value_counts())  
      print(df1.notRepairedDamage.value_counts())
```

```
manuell      227724  
automatik    64153  
Name: gearbox, dtype: int64  
test        153806  
control     143319  
Name: abtest, dtype: int64  
nein        228640  
ja          26457  
Name: notRepairedDamage, dtype: int64
```

We decided to keep gearbox and notRepairedDamage since they differ among the dataset and could be useful for model

Although abtest is A/B testing(user experience research) of eBay, so we decided to drop it.

```
[30]: df1 = df1.drop(["abtest"], axis='columns')
```

3.1.5 offerType nrOfPictures seller

```
[31]: print(df1.offerType.value_counts())  
      print(df1.nrOfPictures.value_counts())  
      print(df1.seller.value_counts())
```

```
Angebot      297123  
Gesuch         2  
Name: offerType, dtype: int64  
0      297125  
Name: nrOfPictures, dtype: int64  
privat      297124  
gewerblich    1  
Name: seller, dtype: int64
```

Since there are just 2 unique values and 1-2 value in second ones in seller and offerType we decided to drop them.

If it comes to nrOfPictures looks like there was a problem with scrawler and we have only zeros. That's why we have to drop them since it will not help to us how it is right now and there is no way to predict it at the moment.

```
[32]: df1 = df1.drop(["offerType", "nrOfPictures", "seller"], axis='columns')
```

```
[33]: df1.head(10)
```



```
[33]:
```

	price	vehicleType	gearbox	powerPS	model	kilometer	fuelType	\
1	9.814656	coupe	manuell	5.247024	NaN	125000	diesel	
2	9.190138	suv	automatik	5.093750	grand	125000	diesel	
3	7.313220	kleinwagen	manuell	4.317488	golf	150000	benzin	
4	8.188689	kleinwagen	manuell	4.234107	fabia	90000	diesel	
5	6.476972	limousine	manuell	4.624973	3er	150000	benzin	
6	7.696213	cabrio	manuell	4.691348	2_reihe	150000	benzin	
8	9.581904	bus	manuell	4.828314	c_max	30000	benzin	
9	6.906755	kleinwagen	manuell	4.615121	golf	150000	NaN	
10	7.600902	limousine	manuell	4.653960	3_reihe	150000	benzin	
11	7.937017	kombi	manuell	4.941642	passat	150000	diesel	

	brand	notRepairedDamage	priceNormal	daysBeforeSold	namelen	age
1	audi	ja	18300	14	20	6
2	jeep	NaN	9800	22	30	13
3	volkswagen	nein	1500	0	18	16
4	skoda	nein	3600	6	30	9
5	bmw	ja	650	2	50	22
6	peugeot	nein	2200	4	27	13
8	ford	NaN	14500	0	36	3
9	volkswagen	NaN	999	14	53	19
10	mazda	nein	2000	11	17	13
11	volkswagen	ja	2799	0	45	12

3.1.6 Translating data from German to English

Since we finished feature engineering and selection before moving to dealing with missing values we translate data from German to English.

```
[34]: ## translate data from german to english
df1['gearbox'] = df1['gearbox'].replace(to_replace=['automatik', 'manuell'],
    ↳value=['automatic', 'manual'], inplace=False, limit=None)
df1['fuelType'] = df1['fuelType'].replace(to_replace=['andere', 'benzin',
    ↳'elektro'], value=['others', 'petrol', 'electric'], inplace=False,
    ↳limit=None)
df1['vehicleType'] = df1['vehicleType'].replace(to_replace=['andere', 'kombi',
    ↳'kleinwagen'], value=['others', 'station wagon', 'small car'],
    ↳inplace=False, limit=None)
df1['notRepairedDamage'] = df1['notRepairedDamage'].replace(to_replace=['ja',
    ↳'nein'], value=['yes', 'no'], inplace=False, limit=None)
df1['brand'] = df1['brand'].replace(to_replace = ["sonstige_autos"],
    ↳value=["other"], inplace=False, limit=None)
```

3.2 Missing values

We have removed problematic values like duplicates and outliers, after that we have chosed the features we are going to work with and drop the rest. Now we are going to deal with missing values

in the dataset.

```
[35]: dfdpna = df1.copy()
dfdpna = dfdpna.dropna()

missing = 100 - 100 * dfdpna['price'].count() / df1['price'].count()
print("Missing values:", missing, "%")
```

Missing values: 21.833235170382835 %

It looks like around 22% of the data has missing values.

3.3 Imputation

We decided to investigate every column with missing values one by one and see how we can deal with them.

```
[36]: dfim = df1.copy()
```

```
[37]: dfim.isnull().sum()
```

```
[37]: price                0
vehicleType             17107
gearbox                 5248
powerPS                 0
model                  10686
kilometer               0
fuelType               16998
brand                   0
notRepairedDamage      42028
priceNormal             0
daysBeforeSold         0
namelen                 0
age                     0
dtype: int64
```

3.3.1 vehicleType

vehicleType has 17107 missing values.

This column there are 8 categorical values and their distribution looks like as it follows.

```
[38]: dfim['vehicleType'].value_counts()
```

```
[38]: limousine           80872
small car              65554
station wagon         58854
bus                   26371
cabrio                19283
coupe                 14652
```

```
suv          12267
others       2165
Name: vehicleType, dtype: int64
```

We checked the mean of every group in order to do some meaningful imputation.

```
[39]: dfim.groupby(['vehicleType']).priceNormal.mean()
```

```
[39]: vehicleType
bus          6713.635547
cabrio       8792.936213
coupe        8559.776754
limousine    5987.652265
others       4219.128868
small car    3047.271867
station wagon 5968.634859
suv          12458.418032
Name: priceNormal, dtype: float64
```

After seeing difference in the average prices we decided to create a function for filling missing values in vehicleType by using means

```
[40]: def fillnavt(x):
        if type(x) == str:
            return x
        if x >= 12000:
            return "suv"
        if x < 12000 and x >= 10500:
            return "coupe"
        if x < 10500 and x >= 9000:
            return "cabrio"
        if x < 9000 and x >= 6700:
            return "bus"
        if x < 6700 and x >= 6100:
            return "limousine"
        if x < 6100 and x >= 4500:
            return "station wagon"
        if x < 4500:
            return "small car"
```

```
[41]: dfim['Col2'] = dfim.priceNormal.apply(lambda x: fillnavt(x))
```

```
[42]: dfim['vehicleType'] = dfim['vehicleType'].combine_first(dfim['Col2'])
```

```
[43]: dfim['vehicleType'].value_counts()
```

```
[43]: limousine      81167
      small car    78797
      station wagon 60130
      bus          27313
      cabrio       19600
      coupe        14913
      suv          13040
      others       2165
      Name: vehicleType, dtype: int64
```

3.3.2 gearbox

We decided to drop gearbox since it's hard to predict the car's gear.

```
[44]: dfim = dfim[-dfim['gearbox'].isnull()]
```

```
[45]: dfim['model'].unique().size
```

```
[45]: 248
```

3.3.3 Model

```
[46]: dfim.groupby(['brand', 'model']).priceNormal.mean()
```

```
[46]: brand      model
      alfa_romeo  145      1103.093750
                        147      2366.404472
                        156      1661.785855
                        159      7318.851282
                        andere  6192.350120
                        ...
      volvo      v40      1912.220257
                        v50      6066.584746
                        v60     16028.181818
                        v70      4626.147887
                        xc_reihe 15487.252595
      Name: priceNormal, Length: 294, dtype: float64
```

Since there is 251 model and price differ a lot between models, it is very hard to replace the missings in a sensible way.

To not cause bias we decided to drop them.

```
[47]: dfim = dfim[-dfim['model'].isnull()]
```

```
[48]: dfim.isnull().sum()
```

```
[48]: price           0
      vehicleType     0
      gearbox         0
      powerPS         0
      model           0
      kilometer       0
      fuelType        13217
      brand           0
      notRepairedDamage 36471
      priceNormal     0
      daysBeforeSold  0
      namelen         0
      age             0
      Col2            0
      dtype: int64
```

3.3.4 fuelType

There are 13217 missing values in the fuelType. We want to fill those missing values in a sensible way and in order to do so we check average prices and distributions of it.

```
[49]: dfim['fuelType'].value_counts()
```

```
[49]: petrol      173508
      diesel      90085
      lpg         4279
      cng         467
      hybrid      201
      electric     53
      others       43
      Name: fuelType, dtype: int64
```

```
[50]: dfim.groupby(['fuelType']).priceNormal.mean()
```

```
[50]: fuelType
      cng      4971.732334
      diesel  8708.634967
      electric 9604.188679
      hybrid 11885.194030
      lpg    4301.146997
      others  4184.534884
      petrol  4697.426032
      Name: priceNormal, dtype: float64
```

Looks like fuelType is mostly petrol and after that diesel. Also diesel cars are more expensive than cars with petrol.

We decided to write the following function in order to guess the fuelType according to price.

```
[51]: def fillft(x):
      if x >= 6000:
          return "diesel"
      if x < 6000:
          return "petrol"
```

```
[52]: dfim['Col13'] = dfim.priceNormal.apply(lambda x: fillft(x))

dfim['fuelType'] = dfim['fuelType'].combine_first(dfim['Col13'])

dfim['fuelType'].value_counts()
```

```
[52]: petrol      184759
      diesel      92051
      lpg         4279
      cng         467
      hybrid      201
      electric     53
      others       43
      Name: fuelType, dtype: int64
```

```
[53]: dfim.groupby(['fuelType']).priceNormal.mean()
```

```
[53]: fuelType
      cng         4971.732334
      diesel      8758.919338
      electric    9604.188679
      hybrid     11885.194030
      lpg         4301.146997
      others      4184.534884
      petrol      4522.545733
      Name: priceNormal, dtype: float64
```

3.3.5 notRepairedDamage

Missing value in notRepairedDamage highly possible means that, car doesn't have a damage which needs to be repaired

So we decided to fill them with "No"

```
[54]: dfim['notRepairedDamage'].value_counts()
```

```
[54]: no      220650
      yes     24732
      Name: notRepairedDamage, dtype: int64
```

```
[55]: dfim['notRepairedDamage'].fillna(value='no', inplace=True)
```

```
[56]: dfim['notRepairedDamage'].value_counts()
```

```
[56]: no      257121
      yes      24732
      Name: notRepairedDamage, dtype: int64
```

```
[57]: dfim = dfim.drop(["Col2", "Col3"], axis='columns')
```

```
[58]: dfim.sample(10)
```

```
[58]:      price vehicleType  gearbox  powerPS  model  kilometer  \
241901  7.313220  limousine  manual  4.836282  6_reihe  150000
292850  9.433484  limousine  manual  4.465908    golf   40000
61419   7.801391    coupe  manual  5.262690    clk   150000
107362  8.902456  limousine  manual  5.318120    3er   150000
172310  6.318968  small car  manual  4.094345  corsa   150000
294588  8.174703  limousine  manual  4.290459    c3   125000
345536  8.006368  limousine  manual  4.744932    3er   150000
162853  7.762171  small car  manual  4.912655  3_reihe  150000
298023  6.801283  limousine  manual  4.499810  1_reihe  150000
281755  9.200290  small car  automatic  4.653960    golf   30000

      fuelType      brand notRepairedDamage  priceNormal  daysBeforeSold  \
241901  petrol      mazda                yes          1500             15
292850  petrol  volkswagen                no         12500              3
61419   petrol  mercedes_benz                no          2444             21
107362  diesel      bmw                  no          7350              9
172310  petrol      opel                  no           555              4
294588  petrol    citroen                no          3550             31
345536  petrol      bmw                  no           3000              0
162853  petrol    peugeot                no          2350             20
298023  petrol      mazda                no           899              3
281755  petrol  volkswagen                no          9900              4

      namelen  age
241901     26   16
292850     48    4
61419     23   19
107362     16   14
172310     16   21
294588     19   14
345536     12   17
162853     35    0
298023     27   20
281755     18    6
```

```
[59]: dfim.isnull().sum()
```

```
[59]: price          0
      vehicleType    0
      gearbox        0
      powerPS        0
      model          0
      kilometer      0
      fuelType       0
      brand          0
      notRepairedDamage 0
      priceNormal     0
      daysBeforeSold  0
      namelen        0
      age            0
      dtype: int64
```

Since we impute or remove all the missing values successfully we don't need priceNormal in our dataset anymore.

```
[60]: priceNormal = dfim.priceNormal
      dfim = dfim.drop(["priceNormal"], axis='columns')
```

Original dataset with outliers with nas

```
[61]: df.price.size
```

```
[61]: 371528
```

Dataset no outliers with nas

```
[62]: dfl.price.size
```

```
[62]: 297125
```

Original dataset no outliers with imputation of nas

```
[63]: dfim.price.size
```

```
[63]: 281853
```

Original dataset no outliers no nas (drop)

```
[64]: dfdpna.price.size
```

```
[64]: 232253
```

Looks like thanks to imputation we saved 49600 rows(13% of the data)

We will see how it will effect our models


```
[65]: print("Outliers and duplicates(dropped):", 100 - left, "%")
missingdp = 100 - 100 * dfim['price'].count() / df1['price'].count()
print("Missing values(dropped):", missingdp, "%")

print("Missing values(imputed):", missing - missingdp, "%")

leftraw2 = 100 * dfim['price'].count() / df['price'].count()
print("The amount of data left(from raw data):", leftraw2, "%")
```

```
Outliers and duplicates(dropped): 20.026216059085726 %
Missing values(dropped): 5.139924274295325 %
Missing values(imputed): 16.69331089608751 %
The amount of data left(from raw data): 75.86319200706272 %
```

3.3.6 Converting string into integers

We need to: * create ordered numerical labels for notRepairedDamage * do one-hot encoding (binarization) for vehicleType gearbox model fuelType and brand

```
[66]: dfim.dtypes
```

```
[66]: price                float64
vehicleType              object
gearbox                  object
powerPS                  float64
model                    object
kilometer                 int64
fuelType                  object
brand                     object
notRepairedDamage        object
daysBeforeSold           int64
namelen                   int64
age                       int64
dtype: object
```

```
[67]: dfim.head(10)
```

```
[67]:
```

	price	vehicleType	gearbox	powerPS	model	kilometer	fuelType	\
2	9.190138	suv	automatic	5.093750	grand	125000	diesel	
3	7.313220	small car	manual	4.317488	golf	150000	petrol	
4	8.188689	small car	manual	4.234107	fabia	90000	diesel	
5	6.476972	limousine	manual	4.624973	3er	150000	petrol	
6	7.696213	cabrio	manual	4.691348	2_reihe	150000	petrol	
8	9.581904	bus	manual	4.828314	c_max	30000	petrol	
9	6.906755	small car	manual	4.615121	golf	150000	petrol	
10	7.600902	limousine	manual	4.653960	3_reihe	150000	petrol	
11	7.937017	station wagon	manual	4.941642	passat	150000	diesel	
12	6.906755	station wagon	manual	4.744932	passat	150000	petrol	

	brand	notRepairedDamage	daysBeforeSold	namelen	age
2	jeep	no	22	30	13
3	volkswagen	no	0	18	16
4	skoda	no	6	30	9
5	bmw	yes	2	50	22
6	peugeot	no	4	27	13
8	ford	no	0	36	3
9	volkswagen	no	14	53	19
10	mazda	no	11	17	13
11	volkswagen	yes	0	45	12
12	volkswagen	no	17	33	22

We create ordered numerical labels for notRepairedDamage.

```
[68]: nrdDict = {'no':0, "yes":1}
dfim["notRepairedDamage"]=dfim.notRepairedDamage.map(nrdDict)
dfdpna["notRepairedDamage"]=dfdpna.notRepairedDamage.map(nrdDict)
```

We did one-hot encoding for all variables with nominal levels to binary form.

```
[69]: levCols = []
numCols = []
for col in dfim.columns:
    if dfim[col].dtype==object:
        levCols.append(col)
    else:
        numCols.append(col)

[70]: dummimint = pd.get_dummies(dfim[levCols])
a = dfim.drop(['vehicleType', 'gearbox', 'model', 'fuelType', 'brand'],
    ↪axis='columns')
dfimdm = pd.concat([a, dummimint], axis=1, sort=False)

dummdpnaint = pd.get_dummies(dfdpna[levCols])
b = dfdpna.drop(['vehicleType', 'gearbox', 'model', 'fuelType', 'brand'],
    ↪axis='columns')
dfdpnadm = pd.concat([b, dummdpnaint], axis=1, sort=False)
```

Since model has 251 unique variable and it is not possible order them, after doing the one-hot encoding we end up with 310 columns.

```
[71]: dfimdm.head()
```

```
[71]:   price  powerPS  kilometer  notRepairedDamage  daysBeforeSold  namelen  \
2  9.190138  5.093750   125000             0             22         30
3  7.313220  4.317488   150000             0              0         18
4  8.188689  4.234107    90000             0              6         30
```

5	6.476972	4.624973	150000	1	2	50
6	7.696213	4.691348	150000	0	4	27

	age	vehicleType_bus	vehicleType_cabrio	vehicleType_coupe	...	\
2	13	0	0	0	...	
3	16	0	0	0	...	
4	9	0	0	0	...	
5	22	0	0	0	...	
6	13	0	1	0	...	

	brand_saab	brand_seat	brand_skoda	brand_smart	brand_subaru	\
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	1	0	0	
5	0	0	0	0	0	
6	0	0	0	0	0	

	brand_suzuki	brand_toyota	brand_trabant	brand_volkswagen	brand_volvo
2	0	0	0	0	0
3	0	0	0	1	0
4	0	0	0	0	0
5	0	0	0	0	0
6	0	0	0	0	0

[5 rows x 310 columns]

Same steps but without model in order to see the effects.

```
[72]: dfimNoModel = dfim.copy()
dfimNoModel = dfimNoModel.drop(['model'], axis = 'columns')

dfimNoModeldm = pd.get_dummies(dfim[['vehicleType', 'gearbox', 'fuelType', 'brand']])

a = dfimNoModel.drop(['vehicleType', 'gearbox', 'fuelType', 'brand'], axis='columns')

dmNoModel = pd.concat([a, dfimNoModeldm], axis=1, sort=False)
```

After dropping 'model' from our data and doing one-hot encoding we end up with 63 columns. We are thinking that it might be easier to work with especially with KNN method.

```
[73]: dmNoModel.head()
```

	price	powerPS	kilometer	notRepairedDamage	daysBeforeSold	namelen	\
2	9.190138	5.093750	125000	0	22	30	
3	7.313220	4.317488	150000	0	0	18	

4	8.188689	4.234107	90000	0	6	30
5	6.476972	4.624973	150000	1	2	50
6	7.696213	4.691348	150000	0	4	27

	age	vehicleType_bus	vehicleType_cabrio	vehicleType_coupe	...	\
2	13	0	0	0	...	
3	16	0	0	0	...	
4	9	0	0	0	...	
5	22	0	0	0	...	
6	13	0	1	0	...	

	brand_saab	brand_seat	brand_skoda	brand_smart	brand_subaru	\
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	1	0	0	
5	0	0	0	0	0	
6	0	0	0	0	0	

	brand_suzuki	brand_toyota	brand_trabant	brand_volkswagen	brand_volvo
2	0	0	0	0	0
3	0	0	0	1	0
4	0	0	0	0	0
5	0	0	0	0	0
6	0	0	0	0	0

[5 rows x 63 columns]

3.4 Correlation

As we can see in the correlation plot while powerPS, daysBeforeSold and namelen has positive effect on price, kilometer and age has negative effect as we suspected.

Kilometer and age has also positive correlation but that's not a surprise.

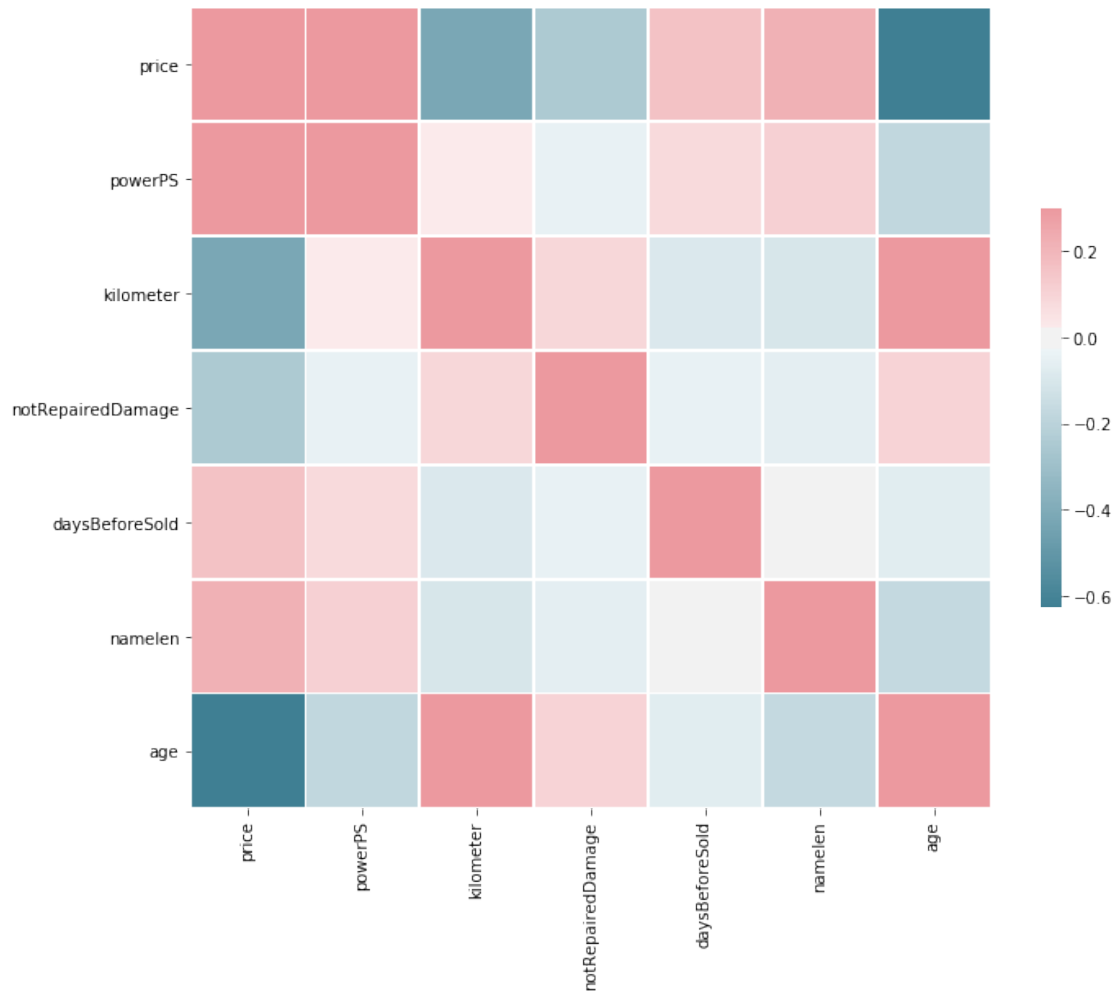
```
[74]: # Compute the correlation matrix
corr = dfim.corr()

# Set up the matplotlib figure
f, ax = plt.subplots(figsize=(11, 9))

# Generate a custom diverging colormap
cmap = sns.diverging_palette(220, 10, as_cmap=True)

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, cmap=cmap, vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})
```

```
[74]: <matplotlib.axes._subplots.AxesSubplot at 0x1ba23c752e8>
```



3.5 Linear regression

We decided to start with linear regression in order to have a base to compare other methods.

First we wanted to check OLS Regression Results to see how every independent variable effecting price.

```
[75]: mod = smf.ols(formula='price ~ vehicleType + gearbox + powerPS + kilometer +
    ↳ fuelType + model + brand + notRepairedDamage', data=dfim)
res = mod.fit()
res.summary()
```

```
[75]: <class 'statsmodels.iolib.summary.Summary'>
      """
```

```

                                OLS Regression Results
=====
Dep. Variable:                  price    R-squared:                  0.700
```

```

Model:                      OLS      Adj. R-squared:      0.700
Method:                    Least Squares      F-statistic:      2188.
Date:                      Wed, 10 Jun 2020      Prob (F-statistic):      0.00
Time:                      20:16:15      Log-Likelihood:      -2.5839e+05
No. Observations:          281853      AIC:      5.174e+05
Df Residuals:              281551      BIC:      5.206e+05
Df Model:                  301
Covariance Type:           nonrobust

```

```

=====
=====

```

	coef	std err	t	P> t
[0.025 0.975]				

Intercept	1.8468	0.064	28.686	0.000
1.721 1.973				
vehicleType[T.cabrio]	0.1488	0.009	16.918	0.000
0.132 0.166				
vehicleType[T.coupe]	-0.1174	0.009	-12.969	0.000
-0.135 -0.100				
vehicleType[T.limousine]	-0.0822	0.007	-11.250	0.000
-0.097 -0.068				
vehicleType[T.others]	-0.2043	0.015	-13.857	0.000
-0.233 -0.175				
vehicleType[T.small car]	-0.2414	0.008	-31.538	0.000
-0.256 -0.226				
vehicleType[T.station wagon]	-0.0899	0.007	-12.197	0.000
-0.104 -0.075				
vehicleType[T.suv]	0.0374	0.012	3.194	0.001
0.014 0.060				
gearbox[T.manual]	-0.0569	0.003	-16.577	0.000
-0.064 -0.050				
fuelType[T.diesel]	0.1762	0.029	6.178	0.000
0.120 0.232				
fuelType[T.electric]	0.9381	0.089	10.568	0.000
0.764 1.112				
fuelType[T.hybrid]	0.3212	0.053	6.071	0.000
0.218 0.425				
fuelType[T.lpg]	-0.2961	0.030	-9.885	0.000
-0.355 -0.237				
fuelType[T.others]	-0.4119	0.097	-4.235	0.000
-0.602 -0.221				
fuelType[T.petrol]	-0.3833	0.029	-13.440	0.000
-0.439 -0.327				
model[T.145]	0.1046	0.119	0.881	0.379
-0.128 0.337				
model[T.147]	0.6549	0.058	11.245	0.000

0.541	0.769				
model[T.156]		-0.0202	0.058	-0.348	0.728
-0.134	0.094				
model[T.159]		0.9781	0.067	14.514	0.000
0.846	1.110				
model[T.1_reihe]		1.2092	0.049	24.612	0.000
1.113	1.306				
model[T.1er]		1.3812	0.064	21.681	0.000
1.256	1.506				
model[T.200]		1.1303	0.352	3.214	0.001
0.441	1.820				
model[T.2_reihe]		1.1847	0.046	25.553	0.000
1.094	1.276				
model[T.300c]		1.6323	0.076	21.471	0.000
1.483	1.781				
model[T.3_reihe]		0.8219	0.046	17.835	0.000
0.732	0.912				
model[T.3er]		0.8472	0.063	13.431	0.000
0.724	0.971				
model[T.4_reihe]		0.5593	0.052	10.670	0.000
0.457	0.662				
model[T.500]		1.8519	0.050	36.696	0.000
1.753	1.951				
model[T.5_reihe]		1.2407	0.054	22.870	0.000
1.134	1.347				
model[T.5er]		0.7099	0.063	11.203	0.000
0.586	0.834				
model[T.601]		1.4441	0.276	5.238	0.000
0.904	1.984				
model[T.6_reihe]		0.7601	0.049	15.447	0.000
0.664	0.857				
model[T.6er]		1.1729	0.078	14.959	0.000
1.019	1.327				
model[T.7er]		0.4934	0.066	7.439	0.000
0.363	0.623				
model[T.80]		0.2014	0.042	4.759	0.000
0.118	0.284				
model[T.850]		0.3803	0.070	5.452	0.000
0.244	0.517				
model[T.90]		0.2780	0.100	2.785	0.005
0.082	0.474				
model[T.900]		0.8701	0.086	10.065	0.000
0.701	1.040				
model[T.9000]		0.8645	0.137	6.301	0.000
0.596	1.133				
model[T.911]		1.0419	0.079	13.223	0.000
0.887	1.196				

model[T.a1]		1.7708	0.046	38.665	0.000
1.681	1.861				
model[T.a2]		1.8264	0.051	35.721	0.000
1.726	1.927				
model[T.a3]		1.0951	0.039	28.225	0.000
1.019	1.171				
model[T.a4]		0.8239	0.039	21.399	0.000
0.748	0.899				
model[T.a5]		1.3844	0.043	32.027	0.000
1.300	1.469				
model[T.a6]		0.6483	0.039	16.658	0.000
0.572	0.725				
model[T.a8]		0.4554	0.047	9.697	0.000
0.363	0.547				
model[T.a_klasse]		0.9897	0.045	21.776	0.000
0.901	1.079				
model[T.accord]		0.9295	0.067	13.921	0.000
0.799	1.060				
model[T.agila]		1.3170	0.058	22.775	0.000
1.204	1.430				
model[T.alhambra]		0.9981	0.094	10.635	0.000
0.814	1.182				
model[T.almera]		0.6720	0.057	11.888	0.000
0.561	0.783				
model[T.altea]		1.2975	0.094	13.809	0.000
1.113	1.482				
model[T.amarok]		1.0028	0.113	8.841	0.000
0.781	1.225				
model[T.andere]		0.9266	0.042	22.016	0.000
0.844	1.009				
model[T.antara]		0.9805	0.079	12.425	0.000
0.826	1.135				
model[T.arosa]		1.5088	0.091	16.625	0.000
1.331	1.687				
model[T.astra]		0.8834	0.048	18.353	0.000
0.789	0.978				
model[T.auris]		1.1629	0.064	18.202	0.000
1.038	1.288				
model[T.avensis]		0.9226	0.055	16.753	0.000
0.815	1.031				
model[T.aveo]		1.6050	0.078	20.639	0.000
1.453	1.757				
model[T.aygo]		1.6855	0.057	29.339	0.000
1.573	1.798				
model[T.b_klasse]		1.4349	0.049	29.164	0.000
1.338	1.531				
model[T.b_max]		2.0223	0.130	15.580	0.000

1.768	2.277				
model[T.beetle]		1.2838	0.055	23.227	0.000
1.175	1.392				
model[T.berlingo]		1.2449	0.051	24.439	0.000
1.145	1.345				
model[T.bora]		0.7539	0.055	13.635	0.000
0.646	0.862				
model[T.boxster]		0.5675	0.075	7.553	0.000
0.420	0.715				
model[T.bravo]		0.8557	0.060	14.224	0.000
0.738	0.974				
model[T.c1]		1.9147	0.056	34.154	0.000
1.805	2.025				
model[T.c2]		1.7804	0.056	31.705	0.000
1.670	1.890				
model[T.c3]		1.7084	0.054	31.775	0.000
1.603	1.814				
model[T.c4]		1.3733	0.055	25.159	0.000
1.266	1.480				
model[T.c5]		0.7872	0.054	14.533	0.000
0.681	0.893				
model[T.c_klasse]		0.7589	0.045	16.919	0.000
0.671	0.847				
model[T.c_max]		1.6816	0.052	32.082	0.000
1.579	1.784				
model[T.c_reihe]		1.5307	0.068	22.500	0.000
1.397	1.664				
model[T.caddy]		1.3847	0.053	25.932	0.000
1.280	1.489				
model[T.calibra]		0.4456	0.068	6.560	0.000
0.312	0.579				
model[T.captiva]		1.2458	0.068	18.228	0.000
1.112	1.380				
model[T.carisma]		0.4578	0.059	7.724	0.000
0.342	0.574				
model[T.carnival]		0.3182	0.067	4.760	0.000
0.187	0.449				
model[T.cayenne]		0.5196	0.078	6.689	0.000
0.367	0.672				
model[T.cc]		1.4227	0.068	20.819	0.000
1.289	1.557				
model[T.ceed]		1.7403	0.065	26.798	0.000
1.613	1.868				
model[T.charade]		0.0631	0.159	0.396	0.692
-0.249	0.375				
model[T.cherokee]		0.8410	0.078	10.800	0.000
0.688	0.994				

model[T.citigo]	1.9904	0.090	22.070	0.000
1.814 2.167				
model[T.civic]	1.0204	0.060	17.131	0.000
0.904 1.137				
model[T.cl]	0.9466	0.059	15.976	0.000
0.831 1.063				
model[T.clio]	1.1741	0.049	23.768	0.000
1.077 1.271				
model[T.clk]	0.8045	0.047	17.094	0.000
0.712 0.897				
model[T.clubman]	0.5661	0.096	5.903	0.000
0.378 0.754				
model[T.colt]	1.3028	0.054	24.238	0.000
1.197 1.408				
model[T.combo]	1.1012	0.058	18.928	0.000
0.987 1.215				
model[T.cooper]	0.4562	0.086	5.287	0.000
0.287 0.625				
model[T.cordoba]	0.7100	0.094	7.516	0.000
0.525 0.895				
model[T.corolla]	0.9233	0.054	17.013	0.000
0.817 1.030				
model[T.corsa]	1.3150	0.048	27.319	0.000
1.221 1.409				
model[T.cr_reihe]	1.2148	0.067	18.266	0.000
1.084 1.345				
model[T.croma]	0.5627	0.111	5.081	0.000
0.346 0.780				
model[T.crossfire]	1.6196	0.095	17.042	0.000
1.433 1.806				
model[T.cuore]	0.8917	0.087	10.241	0.000
0.721 1.062				
model[T.cx_reihe]	1.1878	0.074	16.001	0.000
1.042 1.333				
model[T.defender]	2.2918	0.242	9.463	0.000
1.817 2.766				
model[T.delta]	1.7674	0.142	12.483	0.000
1.490 2.045				
model[T.discovery]	1.4359	0.237	6.047	0.000
0.971 1.901				
model[T.doblo]	1.1523	0.059	19.672	0.000
1.038 1.267				
model[T.ducato]	0.8436	0.057	14.777	0.000
0.732 0.955				
model[T.duster]	0.8548	0.152	5.623	0.000
0.557 1.153				
model[T.e_klasse]	0.6944	0.045	15.420	0.000

0.606	0.783				
model[T.elefantino]		1.5818	0.261	6.063	0.000
1.070	2.093				
model[T.eos]		1.1719	0.058	20.236	0.000
1.058	1.285				
model[T.escort]		0.5166	0.053	9.663	0.000
0.412	0.621				
model[T.espace]		0.1571	0.055	2.834	0.005
0.048	0.266				
model[T.exeo]		1.3274	0.107	12.426	0.000
1.118	1.537				
model[T.fabia]		1.7701	0.064	27.633	0.000
1.645	1.896				
model[T.fiesta]		1.7904	0.047	37.822	0.000
1.698	1.883				
model[T.focus]		1.3199	0.047	28.011	0.000
1.228	1.412				
model[T.forester]		0.7604	0.107	7.103	0.000
0.551	0.970				
model[T.forfour]		0.9733	0.115	8.495	0.000
0.749	1.198				
model[T.fortwo]		1.1146	0.110	10.143	0.000
0.899	1.330				
model[T.fox]		1.8658	0.056	33.162	0.000
1.756	1.976				
model[T.freelander]		1.3894	0.239	5.812	0.000
0.921	1.858				
model[T.fusion]		1.9567	0.060	32.849	0.000
1.840	2.073				
model[T.g_klasse]		1.7848	0.077	23.283	0.000
1.635	1.935				
model[T.galant]		0.1469	0.066	2.241	0.025
0.018	0.275				
model[T.galaxy]		1.1274	0.051	22.300	0.000
1.028	1.227				
model[T.getz]		1.6072	0.057	28.108	0.000
1.495	1.719				
model[T.gl]		1.3157	0.102	12.896	0.000
1.116	1.516				
model[T.glk]		1.2267	0.062	19.743	0.000
1.105	1.348				
model[T.golf]		1.0616	0.051	20.853	0.000
0.962	1.161				
model[T.grand]		1.0174	0.055	18.586	0.000
0.910	1.125				
model[T.i3]		0.7359	0.255	2.881	0.004
0.235	1.237				

model[T.i_reihe]	1.7304	0.050	34.802	0.000
1.633 1.828				
model[T.ibiza]	1.5529	0.089	17.495	0.000
1.379 1.727				
model[T.impreza]	0.4156	0.101	4.105	0.000
0.217 0.614				
model[T.insignia]	1.2794	0.054	23.816	0.000
1.174 1.385				
model[T.jazz]	1.7915	0.068	26.506	0.000
1.659 1.924				
model[T.jetta]	1.1531	0.064	18.092	0.000
1.028 1.278				
model[T.jimny]	1.6136	0.065	24.867	0.000
1.486 1.741				
model[T.juke]	1.5374	0.079	19.400	0.000
1.382 1.693				
model[T.justy]	0.5270	0.109	4.833	0.000
0.313 0.741				
model[T.ka]	1.4858	0.049	30.596	0.000
1.391 1.581				
model[T.kadett]	0.5534	0.075	7.369	0.000
0.406 0.701				
model[T.kaefer]	2.2031	0.116	19.054	0.000
1.976 2.430				
model[T.kalina]	0.7649	0.334	2.288	0.022
0.110 1.420				
model[T.kalos]	1.5361	0.097	15.854	0.000
1.346 1.726				
model[T.kangoo]	1.0455	0.053	19.828	0.000
0.942 1.149				
model[T.kappa]	0.8499	0.188	4.532	0.000
0.482 1.217				
model[T.kuga]	1.7540	0.057	30.574	0.000
1.642 1.866				
model[T.laguna]	0.2409	0.051	4.681	0.000
0.140 0.342				
model[T.lancer]	1.2222	0.066	18.587	0.000
1.093 1.351				
model[T.lanos]	0.9721	0.102	9.486	0.000
0.771 1.173				
model[T.legacy]	0.1467	0.105	1.390	0.164
-0.060 0.353				
model[T.leon]	1.2703	0.090	14.181	0.000
1.095 1.446				
model[T.lodgy]	0.7458	0.184	4.061	0.000
0.386 1.106				
model[T.logan]	0.8080	0.150	5.399	0.000

0.515	1.101				
model[T.lupo]		1.4643	0.053	27.844	0.000
1.361	1.567				
model[T.lybra]		0.1121	0.130	0.863	0.388
-0.142	0.367				
model[T.m_klasse]		0.8297	0.049	16.765	0.000
0.733	0.927				
model[T.m_reihe]		0.9758	0.080	12.205	0.000
0.819	1.133				
model[T.materia]		1.3600	0.166	8.176	0.000
1.034	1.686				
model[T.matiz]		1.6560	0.059	27.896	0.000
1.540	1.772				
model[T.megane]		0.6688	0.050	13.482	0.000
0.572	0.766				
model[T.meriva]		1.1897	0.052	22.974	0.000
1.088	1.291				
model[T.micra]		1.3602	0.052	26.120	0.000
1.258	1.462				
model[T.mii]		2.0504	0.113	18.160	0.000
1.829	2.272				
model[T.modus]		1.2621	0.061	20.853	0.000
1.143	1.381				
model[T.mondeo]		0.8577	0.048	17.955	0.000
0.764	0.951				
model[T.move]		0.7553	0.131	5.759	0.000
0.498	1.012				
model[T.musa]		1.9493	0.162	12.035	0.000
1.632	2.267				
model[T.mustang]		1.4954	0.071	21.089	0.000
1.356	1.634				
model[T.mx_reihe]		1.1171	0.050	22.205	0.000
1.018	1.216				
model[T.navara]		1.2215	0.078	15.694	0.000
1.069	1.374				
model[T.niva]		1.4055	0.154	9.104	0.000
1.103	1.708				
model[T.note]		1.4948	0.074	20.137	0.000
1.349	1.640				
model[T.nubira]		0.6658	0.109	6.084	0.000
0.451	0.880				
model[T.octavia]		1.3846	0.064	21.601	0.000
1.259	1.510				
model[T.omega]		-0.2049	0.051	-4.005	0.000
-0.305	-0.105				
model[T.one]		0.8781	0.088	10.009	0.000
0.706	1.050				

model[T.outlander]	1.5699	0.075	21.049	0.000
1.424 1.716				
model[T.pajero]	1.1979	0.066	18.098	0.000
1.068 1.328				
model[T.panda]	1.7906	0.052	34.277	0.000
1.688 1.893				
model[T.passat]	0.7819	0.051	15.248	0.000
0.681 0.882				
model[T.phaeton]	0.6811	0.068	10.089	0.000
0.549 0.813				
model[T.picanto]	2.0694	0.062	33.310	0.000
1.948 2.191				
model[T.polo]	1.4358	0.051	28.033	0.000
1.335 1.536				
model[T.primera]	0.4019	0.056	7.137	0.000
0.292 0.512				
model[T.ptcruiser]	1.2115	0.065	18.619	0.000
1.084 1.339				
model[T.punto]	1.2169	0.047	25.919	0.000
1.125 1.309				
model[T.q3]	1.2182	0.056	21.886	0.000
1.109 1.327				
model[T.q5]	1.3957	0.050	27.979	0.000
1.298 1.493				
model[T.q7]	1.1500	0.053	21.586	0.000
1.046 1.254				
model[T.qashqai]	1.5911	0.056	28.415	0.000
1.481 1.701				
model[T.r19]	0.2552	0.082	3.100	0.002
0.094 0.417				
model[T.range_rover]	1.1596	0.246	4.709	0.000
0.677 1.642				
model[T.range_rover_evoque]	1.6421	0.249	6.584	0.000
1.153 2.131				
model[T.range_rover_sport]	1.5029	0.245	6.123	0.000
1.022 1.984				
model[T.rangerover]	1.8623	0.276	6.738	0.000
1.321 2.404				
model[T.rav]	0.9549	0.059	16.133	0.000
0.839 1.071				
model[T.rio]	1.5225	0.062	24.392	0.000
1.400 1.645				
model[T.roadster]	1.1239	0.120	9.351	0.000
0.888 1.359				
model[T.roomster]	1.8212	0.073	24.846	0.000
1.678 1.965				
model[T.rx_reihe]	0.6012	0.073	8.282	0.000

0.459	0.743				
model[T.s60]		1.2958	0.072	18.016	0.000
1.155	1.437				
model[T.s_klasse]		0.6658	0.050	13.317	0.000
0.568	0.764				
model[T.s_max]		1.7983	0.055	32.472	0.000
1.690	1.907				
model[T.s_type]		0.2808	0.083	3.380	0.001
0.118	0.444				
model[T.samara]		0.7614	0.306	2.491	0.013
0.162	1.360				
model[T.sandero]		1.0600	0.150	7.072	0.000
0.766	1.354				
model[T.santa]		1.1583	0.059	19.643	0.000
1.043	1.274				
model[T.scenic]		0.5441	0.051	10.657	0.000
0.444	0.644				
model[T.scirocco]		1.4196	0.058	24.664	0.000
1.307	1.532				
model[T.seicento]		1.0837	0.054	20.242	0.000
0.979	1.189				
model[T.sharan]		0.8614	0.054	16.003	0.000
0.756	0.967				
model[T.signum]		0.6907	0.056	12.363	0.000
0.581	0.800				
model[T.sirion]		1.1011	0.096	11.441	0.000
0.912	1.290				
model[T.sl]		1.1197	0.054	20.546	0.000
1.013	1.227				
model[T.slk]		0.9618	0.048	20.033	0.000
0.868	1.056				
model[T.sorento]		1.2642	0.063	20.115	0.000
1.141	1.387				
model[T.spark]		1.6810	0.074	22.653	0.000
1.536	1.826				
model[T.spider]		0.7522	0.075	10.085	0.000
0.606	0.898				
model[T.sportage]		1.5523	0.064	24.225	0.000
1.427	1.678				
model[T.sprinter]		1.1359	0.051	22.319	0.000
1.036	1.236				
model[T.stilo]		0.6264	0.053	11.786	0.000
0.522	0.731				
model[T.superb]		1.3854	0.070	19.760	0.000
1.248	1.523				
model[T.swift]		1.2914	0.053	24.223	0.000
1.187	1.396				

model[T.terios]	1.1534	0.137	8.427	0.000
0.885	1.422			
model[T.tigra]	0.8146	0.054	15.219	0.000
0.710	0.920			
model[T.tiguan]	1.3716	0.056	24.545	0.000
1.262	1.481			
model[T.toledo]	0.5895	0.096	6.171	0.000
0.402	0.777			
model[T.touareg]	0.8677	0.058	15.084	0.000
0.755	0.980			
model[T.touran]	1.2259	0.052	23.530	0.000
1.124	1.328			
model[T.transit]	1.5779	0.053	29.979	0.000
1.475	1.681			
model[T.transporter]	1.4610	0.052	28.247	0.000
1.360	1.562			
model[T.tt]	1.1289	0.042	26.875	0.000
1.047	1.211			
model[T.tucson]	1.3483	0.063	21.474	0.000
1.225	1.471			
model[T.twingo]	1.1578	0.049	23.563	0.000
1.061	1.254			
model[T.up]	1.9380	0.061	31.829	0.000
1.819	2.057			
model[T.v40]	0.6088	0.057	10.645	0.000
0.497	0.721			
model[T.v50]	1.5457	0.065	23.734	0.000
1.418	1.673			
model[T.v60]	1.7313	0.105	16.488	0.000
1.525	1.937			
model[T.v70]	0.9641	0.058	16.699	0.000
0.851	1.077			
model[T.v_klasse]	0.7042	0.071	9.981	0.000
0.566	0.842			
model[T.vectra]	0.3085	0.049	6.313	0.000
0.213	0.404			
model[T.verso]	1.0629	0.060	17.717	0.000
0.945	1.181			
model[T.viano]	1.3865	0.056	24.837	0.000
1.277	1.496			
model[T.vito]	1.0376	0.050	20.853	0.000
0.940	1.135			
model[T.vivaro]	1.3636	0.061	22.338	0.000
1.244	1.483			
model[T.voyager]	0.8044	0.066	12.246	0.000
0.676	0.933			
model[T.wrangler]	1.5290	0.079	19.313	0.000

1.374	1.684				
model[T.x_reihe]		1.1355	0.065	17.581	0.000
1.009	1.262				
model[T.x_trail]		1.0725	0.067	16.107	0.000
0.942	1.203				
model[T.x_type]		0.7191	0.079	9.113	0.000
0.564	0.874				
model[T.xc_reihe]		1.6091	0.063	25.350	0.000
1.485	1.734				
model[T.yaris]		1.4819	0.052	28.448	0.000
1.380	1.584				
model[T.yeti]		1.5807	0.077	20.581	0.000
1.430	1.731				
model[T.ypsilon]		1.7062	0.096	17.748	0.000
1.518	1.895				
model[T.z_reihe]		1.2700	0.067	19.058	0.000
1.139	1.401				
model[T.zafira]		0.8742	0.049	17.706	0.000
0.777	0.971				
brand[T.audi]		0.1514	0.035	4.347	0.000
0.083	0.220				
brand[T.bmw]		0.0831	0.055	1.499	0.134
-0.026	0.192				
brand[T.chevrolet]		-0.1161	0.038	-3.084	0.002
-0.190	-0.042				
brand[T.chrysler]		-0.7256	0.044	-16.625	0.000
-0.811	-0.640				
brand[T.citroen]		-0.1802	0.034	-5.261	0.000
-0.247	-0.113				
brand[T.dacia]		0.5313	0.142	3.738	0.000
0.253	0.810				
brand[T.daewoo]		-0.6205	0.060	-10.332	0.000
-0.738	-0.503				
brand[T.daihatsu]		0.2599	0.072	3.598	0.000
0.118	0.402				
brand[T.fiat]		-0.0164	0.034	-0.478	0.633
-0.084	0.051				
brand[T.ford]		-0.4807	0.036	-13.535	0.000
-0.550	-0.411				
brand[T.honda]		-0.1289	0.048	-2.667	0.008
-0.224	-0.034				
brand[T.hyundai]		-0.2242	0.036	-6.286	0.000
-0.294	-0.154				
brand[T.jaguar]		0.0285	0.053	0.543	0.587
-0.074	0.131				
brand[T.jeep]		-0.1692	0.051	-3.312	0.001
-0.269	-0.069				

brand[T.kia]	-0.3766	0.042	-8.868	0.000
-0.460 -0.293				
brand[T.lada]	-0.4098	0.139	-2.955	0.003
-0.682 -0.138				
brand[T.lancia]	-0.5988	0.078	-7.687	0.000
-0.751 -0.446				
brand[T.land_rover]	-0.2261	0.234	-0.966	0.334
-0.685 0.233				
brand[T.mazda]	-0.0614	0.035	-1.777	0.076
-0.129 0.006				
brand[T.mercedes_benz]	0.1136	0.033	3.460	0.001
0.049 0.178				
brand[T.mini]	1.1033	0.080	13.828	0.000
0.947 1.260				
brand[T.mitsubishi]	-0.3105	0.038	-8.086	0.000
-0.386 -0.235				
brand[T.nissan]	-0.0224	0.039	-0.571	0.568
-0.099 0.054				
brand[T.opel]	0.0124	0.037	0.331	0.740
-0.061 0.085				
brand[T.peugeot]	-0.0273	0.034	-0.797	0.425
-0.094 0.040				
brand[T.porsche]	0.6413	0.060	10.702	0.000
0.524 0.759				
brand[T.renault]	0.0563	0.038	1.497	0.134
-0.017 0.130				
brand[T.rover]	-0.5899	0.046	-12.789	0.000
-0.680 -0.499				
brand[T.saab]	-0.4309	0.045	-9.641	0.000
-0.518 -0.343				
brand[T.seat]	-0.0635	0.083	-0.768	0.442
-0.226 0.099				
brand[T.skoda]	-0.1085	0.055	-1.972	0.049
-0.216 -0.001				
brand[T.smart]	0.5569	0.105	5.283	0.000
0.350 0.763				
brand[T.subaru]	0.2625	0.086	3.043	0.002
0.093 0.432				
brand[T.suzuki]	0.1446	0.036	4.001	0.000
0.074 0.215				
brand[T.toyota]	0.1394	0.038	3.704	0.000
0.066 0.213				
brand[T.trabant]	0.8057	0.231	3.489	0.000
0.353 1.258				
brand[T.volkswagen]	0.0958	0.041	2.328	0.020
0.015 0.176				
brand[T.volvo]	-0.3695	0.042	-8.727	0.000

-0.452	-0.286				
powerPS		1.4694	0.005	293.154	0.000
1.460	1.479				
kilometer		-1.069e-05	3.31e-08	-322.660	0.000
-1.08e-05	-1.06e-05				
notRepairedDamage		-0.6162	0.004	-151.312	0.000
-0.624	-0.608				

```
=====
Omnibus:                28567.517    Durbin-Watson:                2.000
Prob(Omnibus):           0.000    Jarque-Bera (JB):            61264.757
Skew:                    -0.643    Prob(JB):                     0.00
Kurtosis:                4.887    Cond. No.                    8.01e+07
=====
```

Warnings:

```
[1] Standard Errors assume that the covariance matrix of the errors is correctly
specified.
[2] The condition number is large, 8.01e+07. This might indicate that there are
strong multicollinearity or other numerical problems.
"""
```

First step for linear regression prediction is preparing dependent and independent variables

```
[76]: X = dfimdm.drop(['price'], axis='columns')
      y = dfimdm.price
```

We decided to calculate predicted values and check the distribution of errors(differences between real and predicted values).

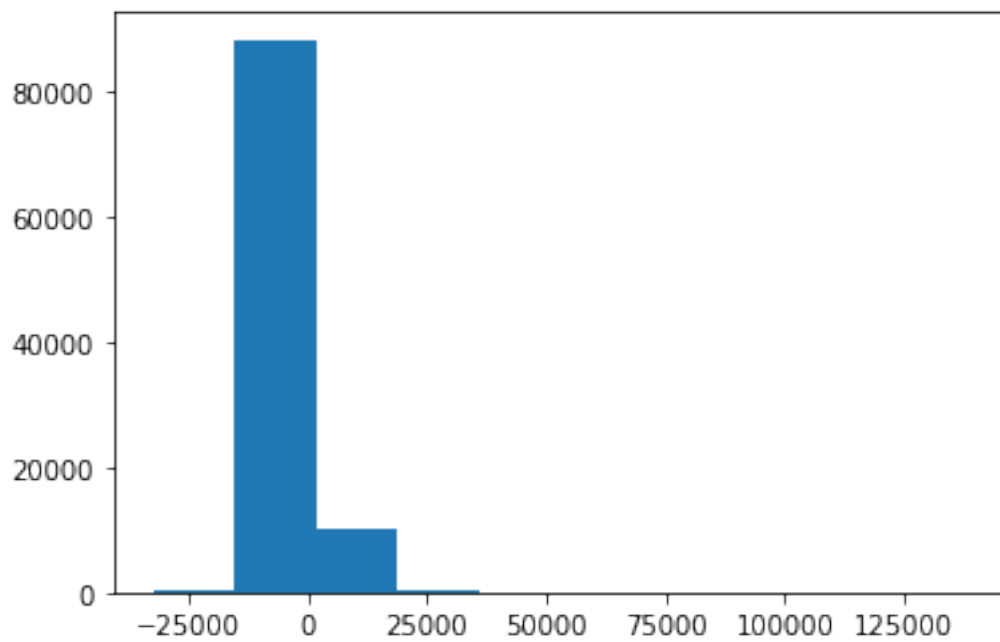
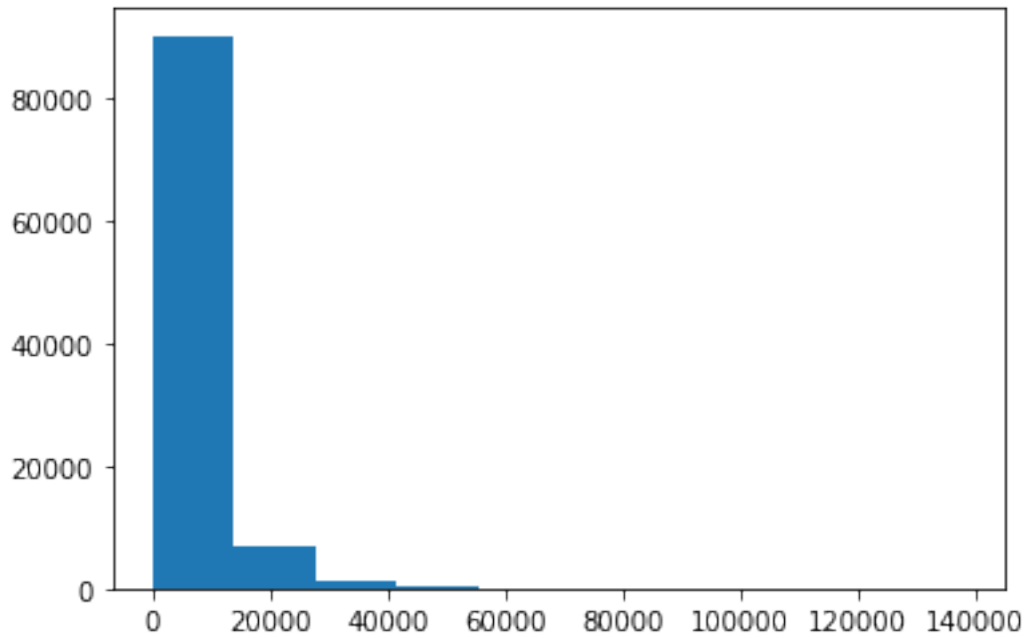
```
[77]: X_train, X_test, y_train, y_testLR = train_test_split(X, y, test_size=0.35)

linreg = LinearRegression()
linreg.fit(X_train, y_train)
y_predLR = linreg.predict(X_test)
scoreLR = linreg.score(X_test, y_testLR)
print("Mean squared error:", np.sqrt(metrics.mean_squared_error(y_testLR,
    ↳y_predLR)))
print("R-squared score:", scoreLR*100, "%")
```

Mean squared error: 0.5129529221824182

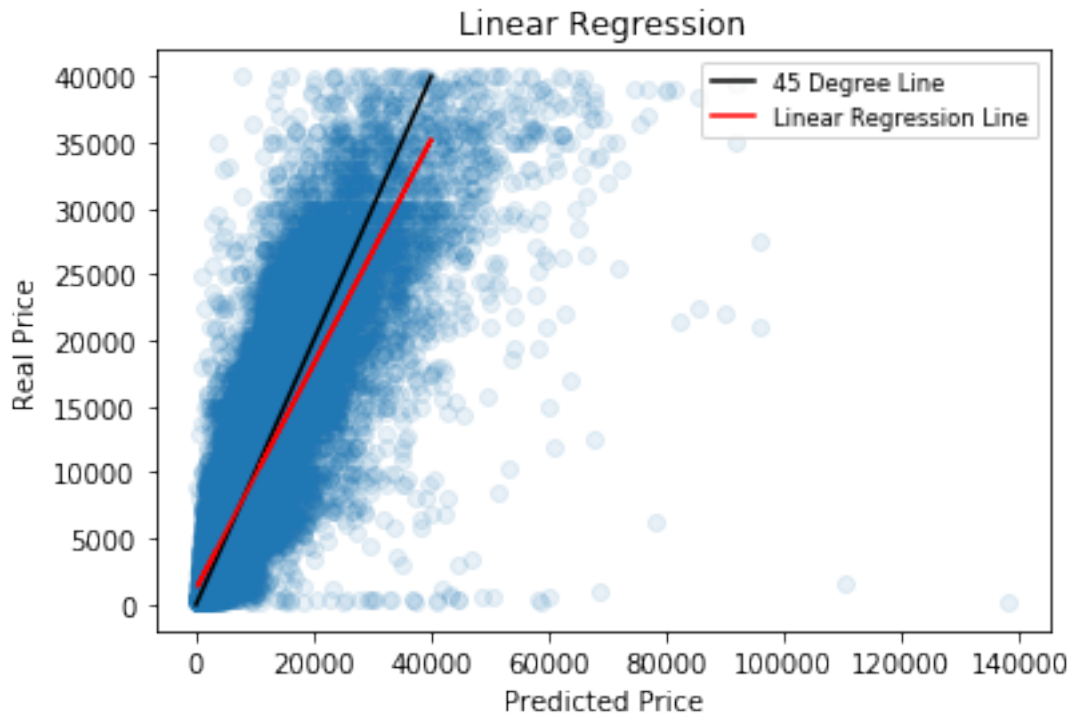
R-squared score: 78.58930520058853 %

```
[78]: plt.hist(np.exp(y_predLR))
      plt.show()
      plt.hist(np.exp(y_predLR) - np.exp(y_testLR.values))
      plt.show()
```



In the first plot we can see how Linear Regression prediction looks like and in the second plot we can see that the difference between predicted and real prices.

```
[79]: plt.scatter(np.exp(y_predLR), np.exp(y_testLR.values), alpha=0.1)
plt.xlabel('Predicted Price')
plt.ylabel('Real Price')
m, b = np.polyfit(np.exp(y_predLR), np.exp(y_testLR.values), 1)
plt.plot([0, 40000], [0, 40000], color = 'black')
plt.plot(np.exp(y_testLR.values), m*np.exp(y_testLR.values) + b, 'red')
plt.legend(["45 Degree Line", "Linear Regression Line"], fontsize = "small")
plt.title("Linear Regression")
plt.show()
```



Looks like with linear regression our estimates was correct 78% but it overshoot for some cars and cause outliers. Since we set 40000 as max in our dataset it should have stay there, but we can clearly see from plot that there are even predictions for 160000. Beside that the line is almost visible.

We decided to try Linear Regression with the data which we removed “model” from it.

```
[80]: X = dmNoModel.drop(['price'], axis='columns')
y = dmNoModel.price

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.35)

linregNM = LinearRegression()
linregNM.fit(X_train, y_train)
```

```

y_predLRnm = linregNM.predict(X_test)
scoreLRnm = linregNM.score(X_test, y_test)
print("Mean squared error:", np.sqrt(metrics.mean_squared_error(y_test,
    ↪y_predLRnm)))
print("R-squared score:", scoreLRnm*100,"%")

```

Mean squared error: 0.5461403062141418

R-squared score: 75.69693562079023 %

It looks like R-squared score decrease 3% we can clearly say that model is very important feature for our models. Having it in our models increases need to computational power and time it takes to run the code. But looks like it worths it.

3.6 K-Nearest Neighbors

Next method we want to use is KNN. It is pretty hard to deal large dataset with KNN but we are thinking it might be successfull.

```
[81]: dfimdm.head()
```

```

[81]:      price  powerPS  kilometer  notRepairedDamage  daysBeforeSold  namelen  \
2  9.190138  5.093750    125000             0             22           30
3  7.313220  4.317488    150000             0              0           18
4  8.188689  4.234107     90000             0              6           30
5  6.476972  4.624973    150000             1              2           50
6  7.696213  4.691348    150000             0              4           27

      age  vehicleType_bus  vehicleType_cabrio  vehicleType_coupe  ...  \
2    13                0                0                0  ...
3    16                0                0                0  ...
4     9                0                0                0  ...
5    22                0                0                0  ...
6    13                0                1                0  ...

      brand_saab  brand_seat  brand_skoda  brand_smart  brand_subaru  \
2              0           0           0           0           0
3              0           0           0           0           0
4              0           0           1           0           0
5              0           0           0           0           0
6              0           0           0           0           0

      brand_suzuki  brand_toyota  brand_trabant  brand_volkswagen  brand_volvo
2              0           0           0           0           0
3              0           0           0           1           0
4              0           0           0           0           0
5              0           0           0           0           0
6              0           0           0           0           0

```

[5 rows x 310 columns]

Firstly we wanted to work on the sample to see how KNN is dealing with our data. If it's worth to trouble.

After taking 50000 rows from our data we split them as 0.7 train and 0.3 test

```
[82]: dfimdmsm = dfimdm.sample(50000)
      # Create train and test set

      train , test = train_test_split(dfimdmsm, test_size = 0.3)

      x_train = train.drop('price', axis=1)
      y_train = train['price']

      x_test = test.drop('price', axis = 1)
      y_test = test['price']
```

Since KNN is a Distance-Based algorithm we need to scale all variables otherwise KNN will not perform optimally.

We decided to scale them between 0 and 1.

```
[83]: # Preprocessing - Scaling the features

      scaler = MinMaxScaler(feature_range=(0, 1))

      x_train_scaled = scaler.fit_transform(x_train)
      x_train = pd.DataFrame(x_train_scaled)

      x_test_scaled = scaler.fit_transform(x_test)
      x_test = pd.DataFrame(x_test_scaled)
```

In order to decide how many neighbors we should have take we decided to check one by one what are the mean squared errors and what are the R-squared score.

```
[84]: # 50000 sample
      # checking the error rate for different n_neighbors

      rmse_val = [] #to store rmse values for different k
      for K in range(10):
          K = K+1
          model = neighbors.KNeighborsRegressor(n_neighbors = K)

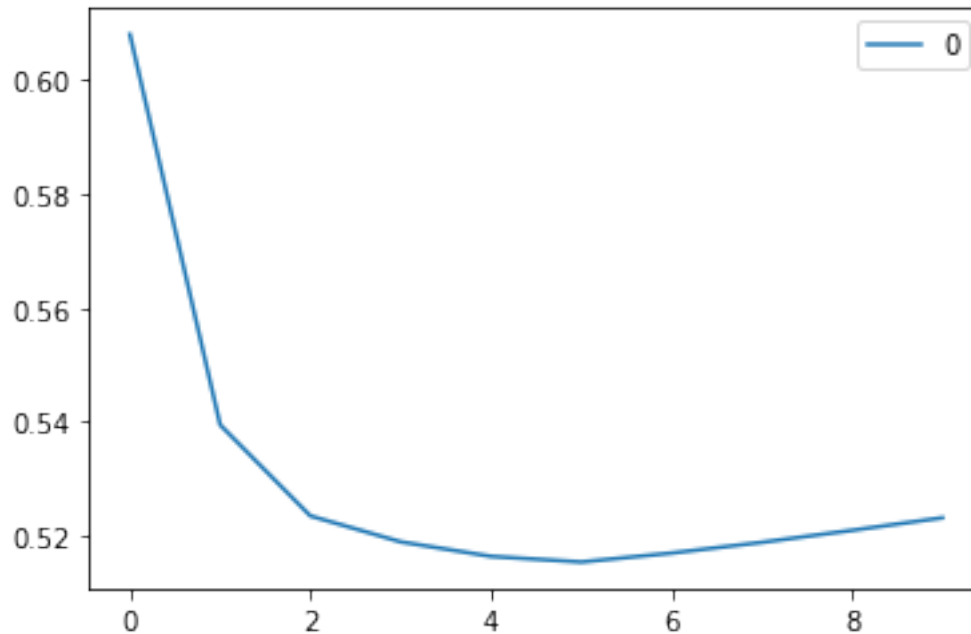
          model.fit(x_train, y_train) #fit the model
          pred=model.predict(x_test) #make prediction on test set
          error = sqrt(mean_squared_error(y_test,pred)) #calculate rmse
          rmse_val.append(error) #store rmse values
```

```
accur = metrics.r2_score(y_test.values, pred)
print('RMSE value for k= ', K , 'is:', error)
print('R-squared score:', accur)
```

```
RMSE value for k= 1 is: 0.6080560421523825
R-squared score: 0.6973588626075502
RMSE value for k= 2 is: 0.5394626768055544
R-squared score: 0.761788046954268
RMSE value for k= 3 is: 0.5234642572165564
R-squared score: 0.7757074701359661
RMSE value for k= 4 is: 0.5189137986664447
R-squared score: 0.7795900568722829
RMSE value for k= 5 is: 0.5163611492951381
R-squared score: 0.7817532117411445
RMSE value for k= 6 is: 0.5153717265257459
R-squared score: 0.7825887953530072
RMSE value for k= 7 is: 0.5169386244209656
R-squared score: 0.7812647839576071
RMSE value for k= 8 is: 0.5188240400932973
R-squared score: 0.7796663006377123
RMSE value for k= 9 is: 0.5209363471358256
R-squared score: 0.7778685433480225
RMSE value for k= 10 is: 0.5231404769646025
R-squared score: 0.7759848492747023
```

```
[85]: # 50000
      # k = 6
      #plotting the rmse values against k values
      curve = pd.DataFrame(rmse_val) #elbow curve
      curve.plot()
```

```
[85]: <matplotlib.axes._subplots.AxesSubplot at 0x1ba3ee8c5c0>
```

```
[86]: kf = KFold(n_splits=5, shuffle=True, random_state=random.randint(0,10000))
errors = []
accurs = []
for train, test in kf.split(dfimdmsm.index.values):
    x = dfimdmsm.drop('price', axis=1)
    y = dfimdmsm['price']
    x_train = x.iloc[train]
    y_train = y.iloc[train]

    x_test = x.iloc[test]
    y_test = y.iloc[test]

    scaler = MinMaxScaler(feature_range=(0, 1))

    x_train_scaled = scaler.fit_transform(x_train)
    x_train = pd.DataFrame(x_train_scaled)

    x_test_scaled = scaler.fit_transform(x_test)
    x_test = pd.DataFrame(x_test_scaled)

    model = neighbors.KNeighborsRegressor(n_neighbors = 6)

    model.fit(x_train, y_train) #fit the model
    pred=model.predict(x_test) #make prediction on test set
    error = sqrt(mean_squared_error(y_test,pred)) #calculate rmse
    errors.append(error) #store rmse values
```

```

    accur = metrics.r2_score(y_test.values, pred)
    accurs.append(accur)
    print('Mean squared error:', error)
    print('R-squared score:', accur)
print('Average Mean squared error:', np.mean(error))
print('Average R-squared score:', np.mean(accur))

```

```

Mean squared error: 0.5097667245894799
R-squared score: 0.7887567176356063
Mean squared error: 0.4972572716109484
R-squared score: 0.7967629504564305
Mean squared error: 0.5130122267239671
R-squared score: 0.7806008658495368
Mean squared error: 0.5073440351737304
R-squared score: 0.788694163689195
Mean squared error: 0.5311165244331162
R-squared score: 0.7682000380565733
Average Mean squared error: 0.5311165244331162
Average R-squared score: 0.7682000380565733

```

It looks like it worth the trouble. Even for 50000 sample it gave same result like linear regression(78%). So we decided to go for it with the all data.

```

[87]: x = dfimdm.drop(['price'], axis='columns')
      y = dfimdm.price

```

Since we are to fit and test with all data this time we were not be able to use Kfold or check what is the best `n_neighbors` but we are guessing that 50 should be enough.

```

[88]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.35)

# Preprocessing - Scaling the features

scaler = MinMaxScaler(feature_range=(0, 1))

x_train_scaled = scaler.fit_transform(x_train)
x_train = pd.DataFrame(x_train_scaled)

x_test_scaled = scaler.fit_transform(x_test)
x_test = pd.DataFrame(x_test_scaled)

# classifier
model = neighbors.KNeighborsRegressor(n_neighbors = 50)

model.fit(x_train, y_train) #fit the model
pred=model.predict(x_test) #make prediction on test set
error = sqrt(mean_squared_error(y_test,pred)) #calculate rmse

```

```

errors.append(error) #store rmse values
accur = metrics.r2_score(y_test.values, pred)
accurs.append(accur)
print('Mean squared error:', error)
print('R-squared score:', accur)

```

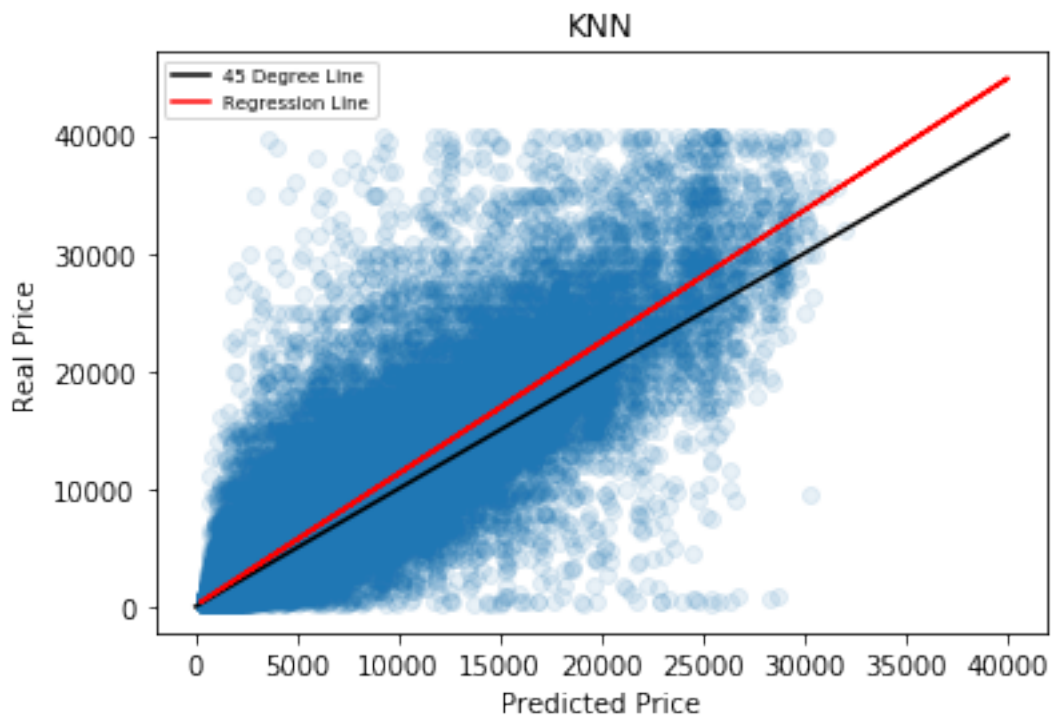
Mean squared error: 0.4991365616561806

R-squared score: 0.7962880626504695

```

[89]: plt.scatter(np.exp(pred), np.exp(y_test.values), alpha=0.1)
plt.xlabel('Predicted Price')
plt.ylabel('Real Price')
m, b = np.polyfit(np.exp(pred), np.exp(y_test.values), 1)
plt.plot([0, 40000], [0, 40000], color = 'black')
plt.plot(np.exp(y_test.values), m*np.exp(y_test.values) + b, 'red')
plt.legend(["45 Degree Line", "Regression Line"], fontsize = "x-small")
plt.title("KNN")
plt.show()

```



We see that KNN with 50 neighbors couldn't predict expensive cars accurately. It did more or less good job until 25k but after that start to not be accurate. Especially after 30k it did not work well at all.

After getting 79% of R-squared score with KNN with 50 neighbors we suspect that maybe same

neighbors number with 50000 sample would work better for all data.

```
[90]: x = dfimdm.drop(['price'], axis='columns')
y = dfimdm.price

x_train, x_test, y_train, y_testknn = train_test_split(x, y, test_size=0.35)

# Preprocessing - Scaling the features

scaler = MinMaxScaler(feature_range=(0, 1))

x_train_scaled = scaler.fit_transform(x_train)
x_train = pd.DataFrame(x_train_scaled)

x_test_scaled = scaler.fit_transform(x_test)
x_test = pd.DataFrame(x_test_scaled)

# classifier
model = neighbors.KNeighborsRegressor(n_neighbors = 6)

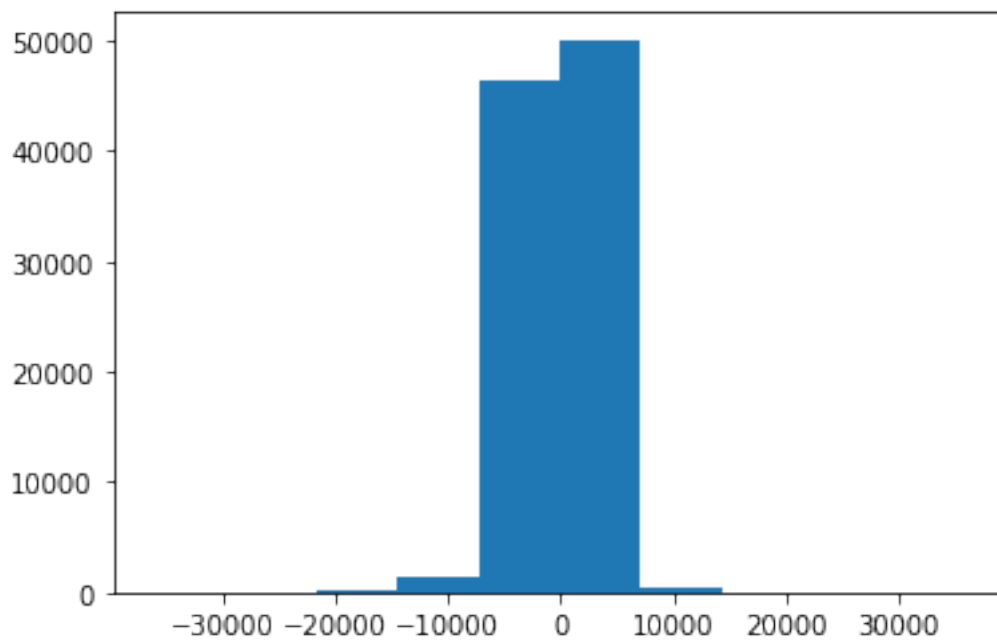
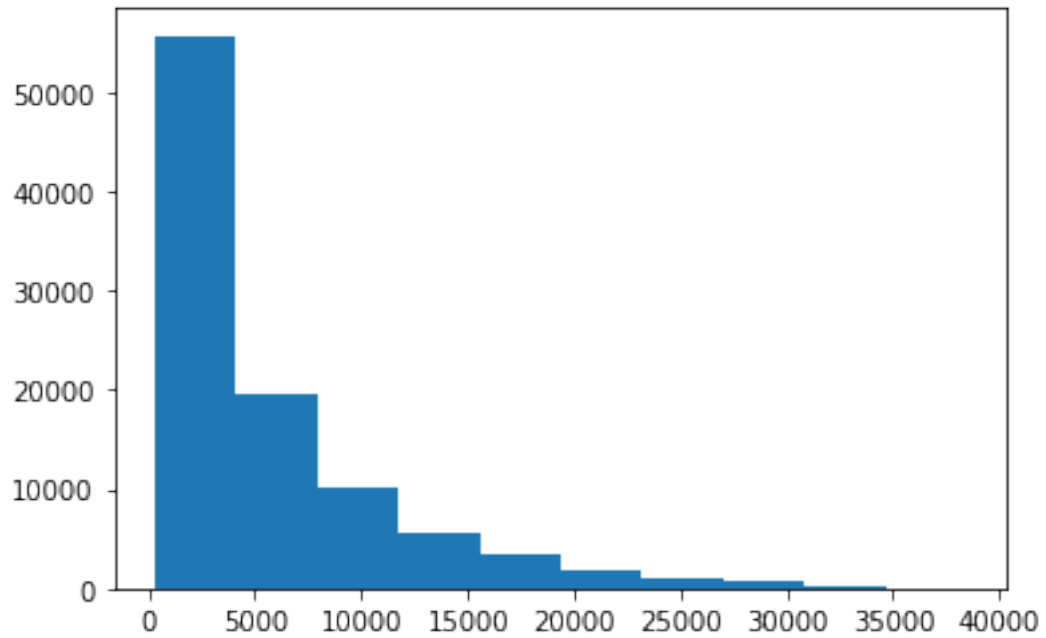
model.fit(x_train, y_train) #fit the model
pred=model.predict(x_test) #make prediction on test set
error = sqrt(mean_squared_error(y_testknn,pred)) #calculate rmse
accur = metrics.r2_score(y_testknn.values, pred)
print('Mean squared error:', error)
print('R-squared score:', accur*100,"%")
```

Mean squared error: 0.44738439622778964

R-squared score: 83.6361413985903 %

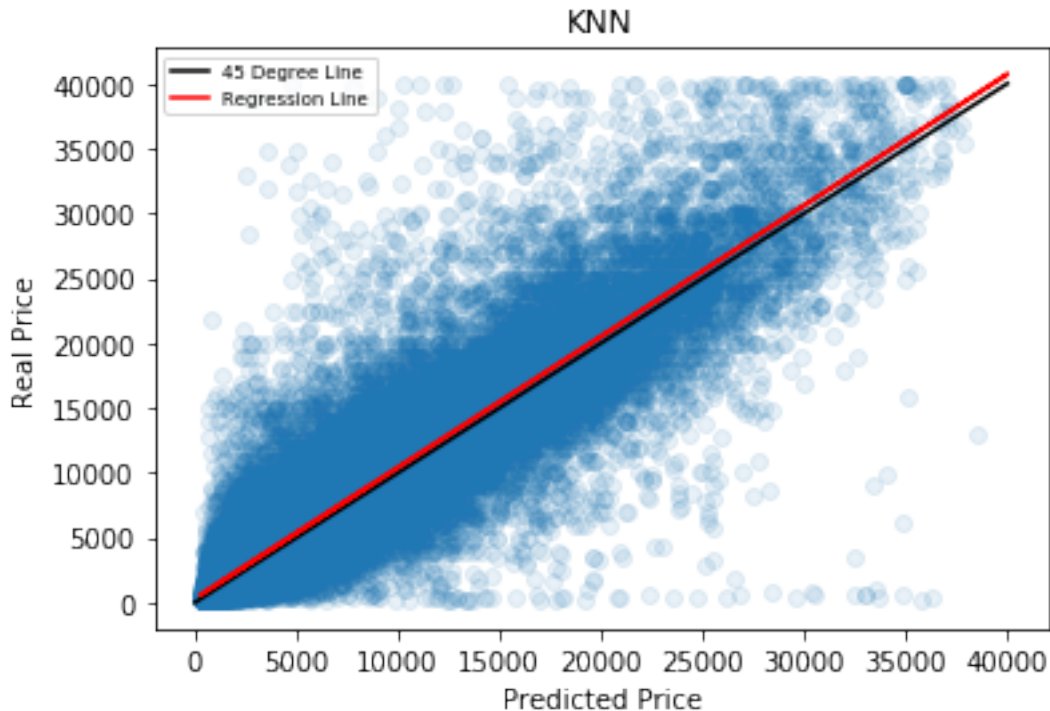
Decreasing numbers of neighbors worked well and it increased R-squared score from 79% to 83%. We wanted to plot the predictions and check how KNN worked.

```
[91]: plt.hist(np.exp(pred))
plt.show()
plt.hist(np.exp(pred) - np.exp(y_testknn))
plt.show()
```



In the first plot we can see how KNN prediction looks like and in the second plot we can see that the difference between predicted and real prices.

```
[92]: plt.scatter(np.exp(pred), np.exp(y_testknn.values), alpha=0.1)
plt.xlabel('Predicted Price')
plt.ylabel('Real Price')
m, b = np.polyfit(np.exp(pred), np.exp(y_testknn.values), 1)
plt.plot([0, 40000], [0, 40000], color = 'black')
plt.plot(np.exp(y_testknn.values), m*np.exp(y_testknn.values) + b, 'red')
plt.legend(["45 Degree Line", "Regression Line"], fontsize = "x-small")
plt.title("KNN")
plt.show()
```



This time looks like predictions became more accurate. Although we can see in general KNN under-price cars. Despite that it's very clear that it made better job than Linear Regression.

Before moving to the next method we want to check again how data without “model” would work with KNN.

```
[93]: x = dmNoModel.drop(['price'], axis='columns')
y = dmNoModel.price
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.35)

# Preprocessing - Scaling the features

scaler = MinMaxScaler(feature_range=(0, 1))
```

```

x_train_scaled = scaler.fit_transform(x_train)
x_train = pd.DataFrame(x_train_scaled)

x_test_scaled = scaler.fit_transform(x_test)
x_test = pd.DataFrame(x_test_scaled)

# classifier
modelNM = neighbors.KNeighborsRegressor(n_neighbors = 6)

modelNM.fit(x_train, y_train) #fit the model
predNM=modelNM.predict(x_test) #make prediction on test set
errorNM = sqrt(mean_squared_error(y_test,predNM)) #calculate rmse
accurNM = metrics.r2_score(y_test.values, predNM)
print('Mean squared error:', errorNM)
print('R-squared score:', accurNM)

```

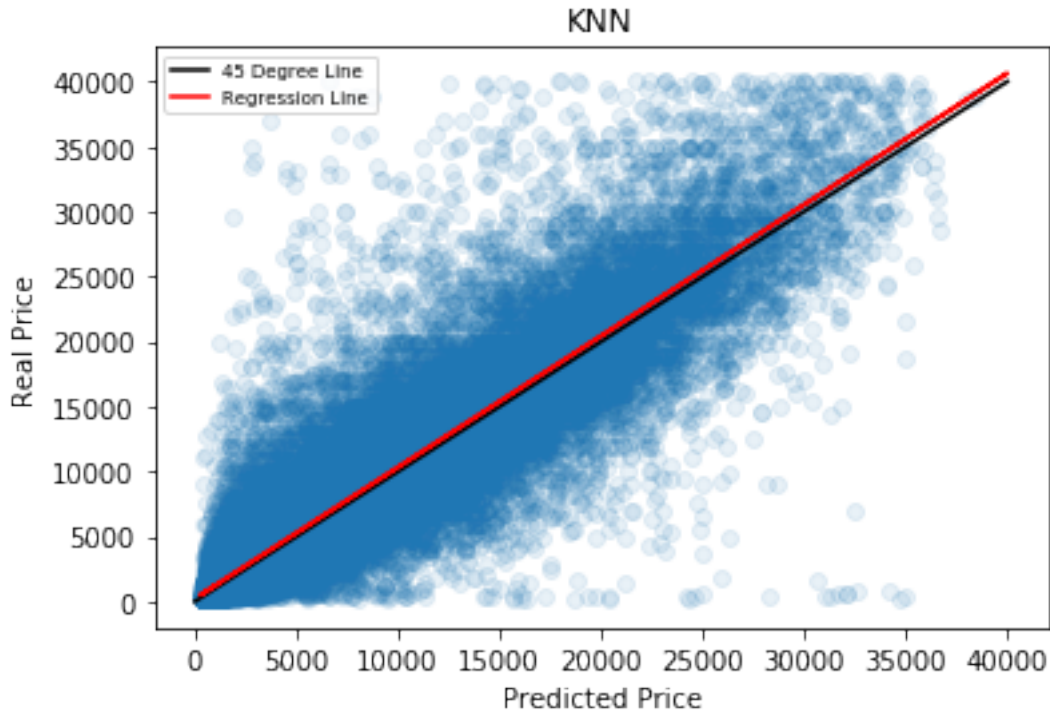
Mean squared error: 0.44659729097938095

R-squared score: 0.8360531864475668

```

[94]: plt.scatter(np.exp(predNM), np.exp(y_test.values), alpha=0.1)
plt.xlabel('Predicted Price')
plt.ylabel('Real Price')
m, b = np.polyfit(np.exp(predNM), np.exp(y_test.values), 1)
plt.plot([0, 40000], [0, 40000], color = 'black')
plt.plot(np.exp(y_test.values), m*np.exp(y_test.values) + b, 'red')
plt.legend(["45 Degree Line", "Regression Line"], fontsize = "x-small")
plt.title("KNN")
plt.show()

```



Unlike Linear Regression in KNN data without “model” worked similar. Since KNN working with distance we are thinking “model” was causing overfitting.

4 Random Forest

Next method we used is random forest. Since we have plenty categories especially under brand and model we are thinking that random forest method will be great fit to our data.

Cross Validation (K-Fold) We split our data into 5 parts, with KFold. And iterate it 5 times, every time test and train data was changing as follows:

```
test-train-train-train-train
train-test-train-train-train
train-train-test-train-train
train-train-train-test-train
train-train-train-train-test
```

We used this method to be able get more from our data and see if R-squared score changes depending on the which part of data we are training and testing our model.

Since random forest is working with tree-based model we wanted to change the way we encode our data and try to make depth shorter.

```
[95]: dfimrf = dfim.copy()
```



```
[96]: for col in ['vehicleType', 'gearbox', 'model', 'fuelType', 'brand']:
        le = preprocessing.LabelEncoder()
        le.fit(dfimrf[col].unique())
        dfimrf[col] = le.transform(dfimrf[col])
```

```
[97]: dfimrf.head()
```

```
[97]:
```

	price	vehicleType	gearbox	powerPS	model	kilometer	fuelType	\
2	9.190138	7	0	5.093750	118	125000	1	
3	7.313220	5	1	4.317488	117	150000	6	
4	8.188689	5	1	4.234107	102	90000	1	
5	6.476972	3	1	4.624973	11	150000	6	
6	7.696213	1	1	4.691348	8	150000	6	

	brand	notRepairedDamage	daysBeforeSold	namelen	age
2	14	0	22	30	13
3	37	0	0	18	16
4	31	0	6	30	9
5	2	1	2	50	22
6	25	0	4	27	13

```
[98]: # Create train and test set

kf = KFold(n_splits=5, shuffle=True, random_state=random.randint(0,10000))
ttscores = []
tnscores = []
errorsRF = []
k = 0
for train, test in kf.split(dfimrf.index.values):
    k += 1
    x = dfimrf.drop('price', axis=1)
    y = dfimrf['price']
    x_train = x.iloc[train]
    y_train = y.iloc[train]

    x_test = x.iloc[test]
    y_testrf = y.iloc[test]

    # classifier
    rfr = RandomForestRegressor()
    rfr.fit(x_train, y_train)
    tnscore = rfr.score(x_train, y_train)
    tnscores.append(tnscore)
    predRF = rfr.predict(x_test)
    ttscore = rfr.score(x_test, y_testrf)
    ttscores.append(ttscore)
```

```

errorRF = sqrt(mean_squared_error(y_testrf,predRF)) #calculate rmse
errorsRF.append(errorRF)

print("Iteration:", k)
print('Mean squared error:', errorRF)
print("Score on train dataset:", tnscore*100,"%")
print("Score on test dataset:", ttscore*100,"%")
print("-----")
print("Average score on train dataset:", np.mean(tnscores)*100,"%")
print("Average score on test dataset:", np.mean(ttscores)*100,"%")
print("Average mean squared error:", np.mean(errorsRF))
scoreRF = np.mean(ttscores)
predRF = rfr.predict(x_test)

```

```

Iteration: 1
Mean squared error: 0.38130632863129355
Score on train dataset: 98.17644246192957 %
Score on test dataset: 88.09088949726018 %
Iteration: 2
Mean squared error: 0.3828974997674672
Score on train dataset: 98.18472038271105 %
Score on test dataset: 88.07246176960331 %
Iteration: 3
Mean squared error: 0.3849012253238724
Score on train dataset: 98.18044549733833 %
Score on test dataset: 87.79987633126763 %
Iteration: 4
Mean squared error: 0.38119757851744873
Score on train dataset: 98.17228778364473 %
Score on test dataset: 88.12875074459188 %
Iteration: 5
Mean squared error: 0.38310438233143723
Score on train dataset: 98.16512373246707 %
Score on test dataset: 88.02976192026522 %
-----
Average score on train dataset: 98.17580397161817 %
Average score on test dataset: 88.02434805259765 %
Average mean squared error: 0.38268140291430386

```

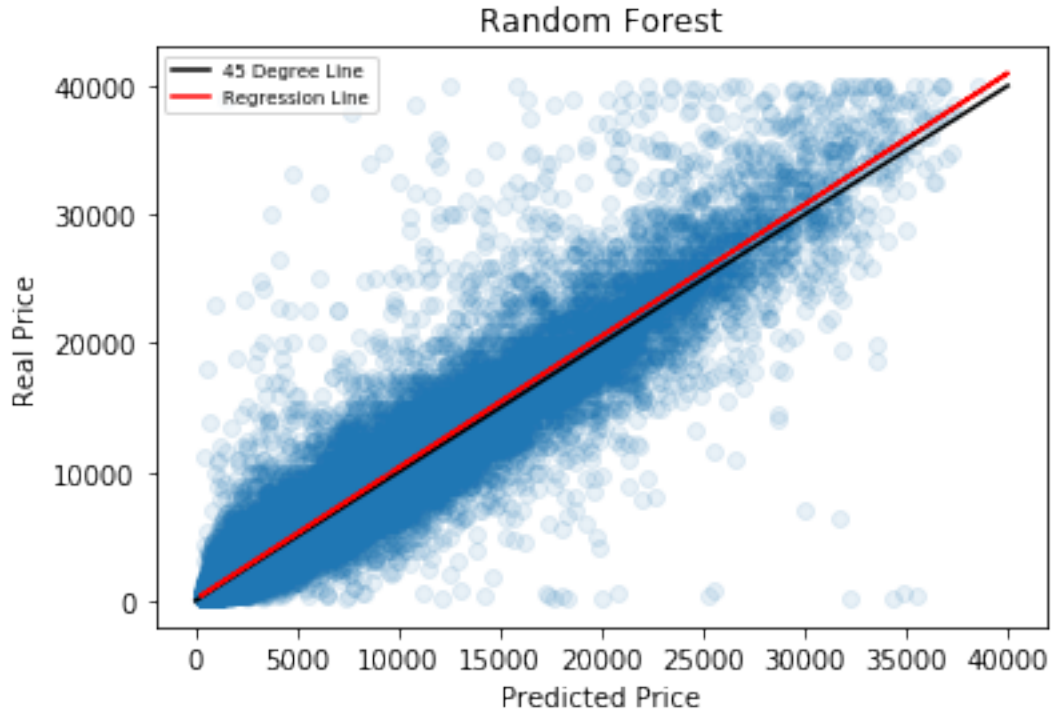
The R-squared score didn't differ much among the data in 5 iteration. Although 88% R-squared score is impressive compare to Linear Regression and KNN methods. Let's see in the plots how predictions of the random forest model looks like.

```

[99]: plt.scatter(np.exp(predRF) ,np.exp(y_testrf), alpha=0.1)
plt.xlabel('Predicted Price')
plt.ylabel('Real Price')
m, b = np.polyfit(np.exp(predRF), np.exp(y_testrf), 1)
plt.plot([0, 40000], [0, 40000], color = 'black')

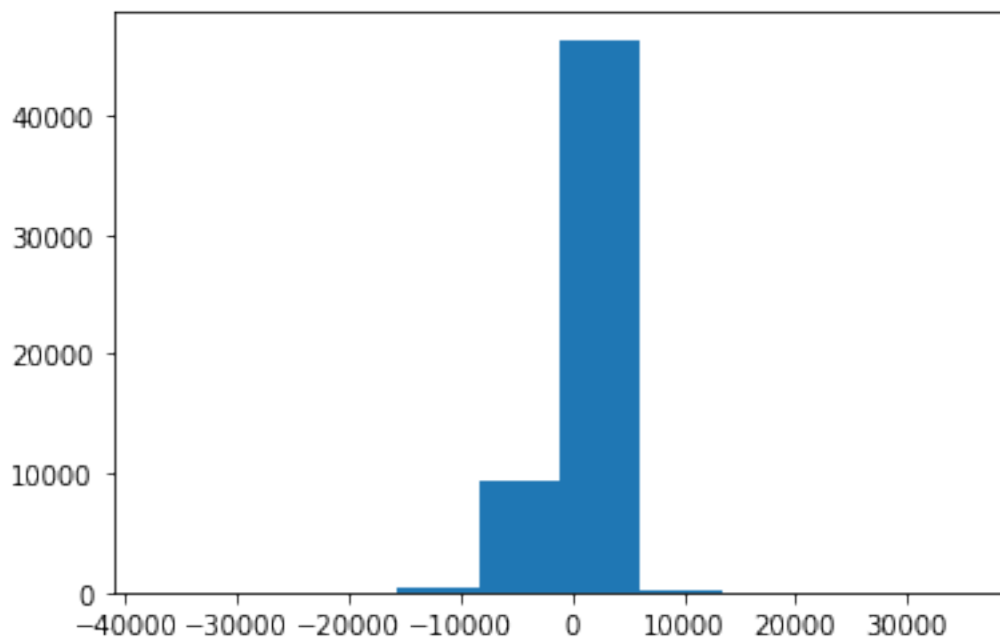
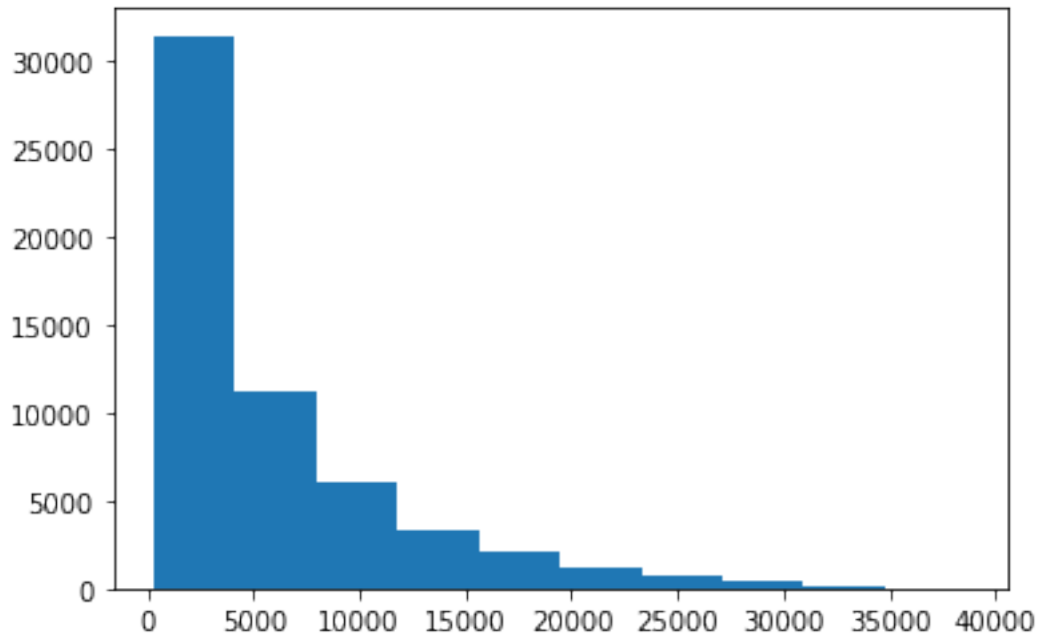
```

```
plt.plot(np.exp(y_testrf.values), m*np.exp(y_testrf.values) + b, 'red')
plt.legend(["45 Degree Line", "Regression Line"], fontsize = "x-small")
plt.title("Random Forest")
plt.show()
```



It's clearly visible that this time model predict better. Of course there are some outliers like over and under pricing but in general the line became thinner and more visible.

```
[100]: plt.hist(np.exp(predRF))
plt.show()
plt.hist(np.exp(predRF) - np.exp(y_testrf.values))
plt.show()
```



In the first plot we can see how Random Forest prediction looks like and in the second plot we can see that the difference between predicted and real prices.

```
[101]: dfimNoModelrf = dfimNoModel.copy()
for col in ['vehicleType', 'gearbox', 'fuelType', 'brand']:
    le = preprocessing.LabelEncoder()
    le.fit(dfimNoModelrf[col].unique())
    dfimNoModelrf[col] = le.transform(dfimNoModelrf[col])

xNM = dfimNoModelrf.drop(['price'], axis='columns')
yNM = dfimNoModelrf.price
```

```
[102]: x_train, x_test, y_train, y_test = train_test_split(xNM, yNM, test_size=0.35)

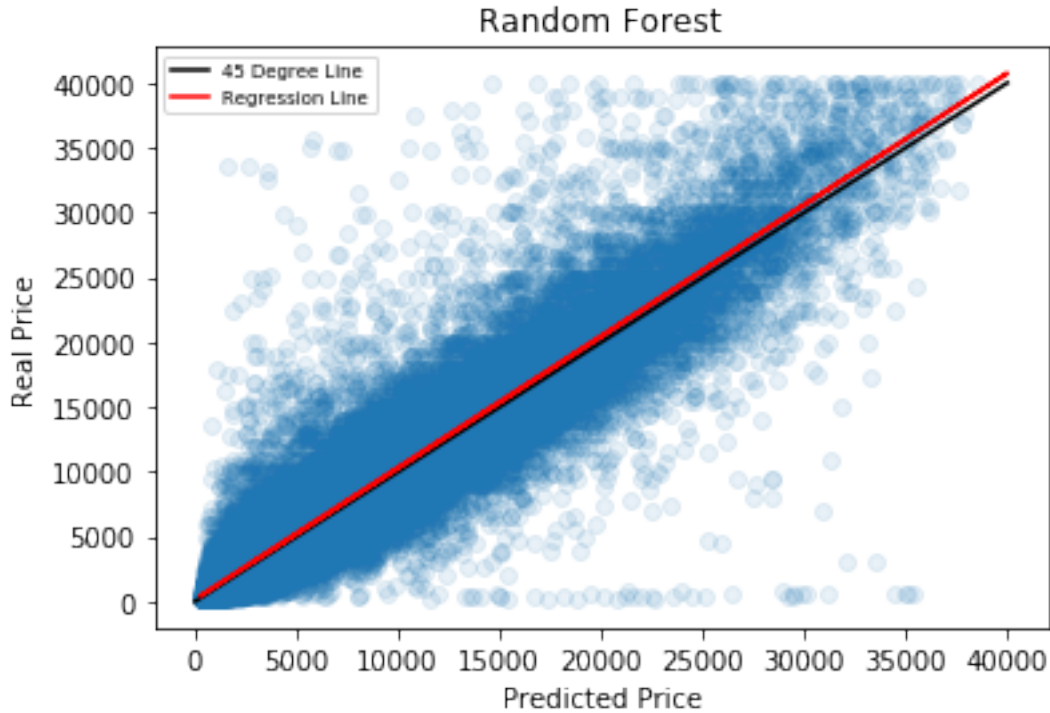
rfrNM = RandomForestRegressor()
rfrNM.fit(x_train, y_train)
predRFNM = rfrNM.predict(x_test)
tnscoreNM = rfrNM.score(x_train, y_train)
ttscoreNM = rfrNM.score(x_test, y_test)

errorRFNM = sqrt(mean_squared_error(y_test, predRFNM))

print('Mean squared error:', errorRFNM)
print("Score on train dataset:", tnscoreNM*100, "%")
print("Score on test dataset:", ttscoreNM*100, "%")
```

Mean squared error: 0.39518677236334177
Score on train dataset: 98.05167813414398 %
Score on test dataset: 87.26288219023658 %

```
[103]: plt.scatter(np.exp(predRFNM), np.exp(y_test.values), alpha=0.1)
plt.xlabel('Predicted Price')
plt.ylabel('Real Price')
m, b = np.polyfit(np.exp(predRFNM), np.exp(y_test.values), 1)
plt.plot([0, 40000], [0, 40000], color = 'black')
plt.plot(np.exp(y_test.values), m*np.exp(y_test.values) + b, 'red')
plt.legend(["45 Degree Line", "Regression Line"], fontsize = "x-small")
plt.title("Random Forest")
plt.show()
```



Despite R-squared score results looks same with and without “model” from plot above we can see that the line was thinner if it comes to predictions with “model”. So we can say that thanks to model in the data our errors are smaller.

4.1 Comparison of Methods and Conclusions

In this project we aimed to find the best regression model for used cars dataset to be able to predict used cars price.

Dataset was scraped with Scrapy from German eBay and there were some mistakes, duplications and outliers. As first step we removed them from the dataset and stayed with 79.9% of the original data.

There were 20 variables in the original dataset and we had to remove “dateCrawled”, “postalCode”, “abtest”, “offerType”, “nrOfPictures”, and “seller” because they were not usefull for our purpose. We created new features: * “daysBeforeSold” by using “dateCreated” and “lastSeen” * “namelen” by using “name” * “age” by using “yearOfRegistration”

After that we worked on missing values columnwise and impute what we could impute meaningfully and drop the rest. At the end we had 75.8% of the original data ready to use for price prediction.

We used log-transformation in price and powerPs and one hot encoding for categorical variables to be able to perform machine learning algorithms.

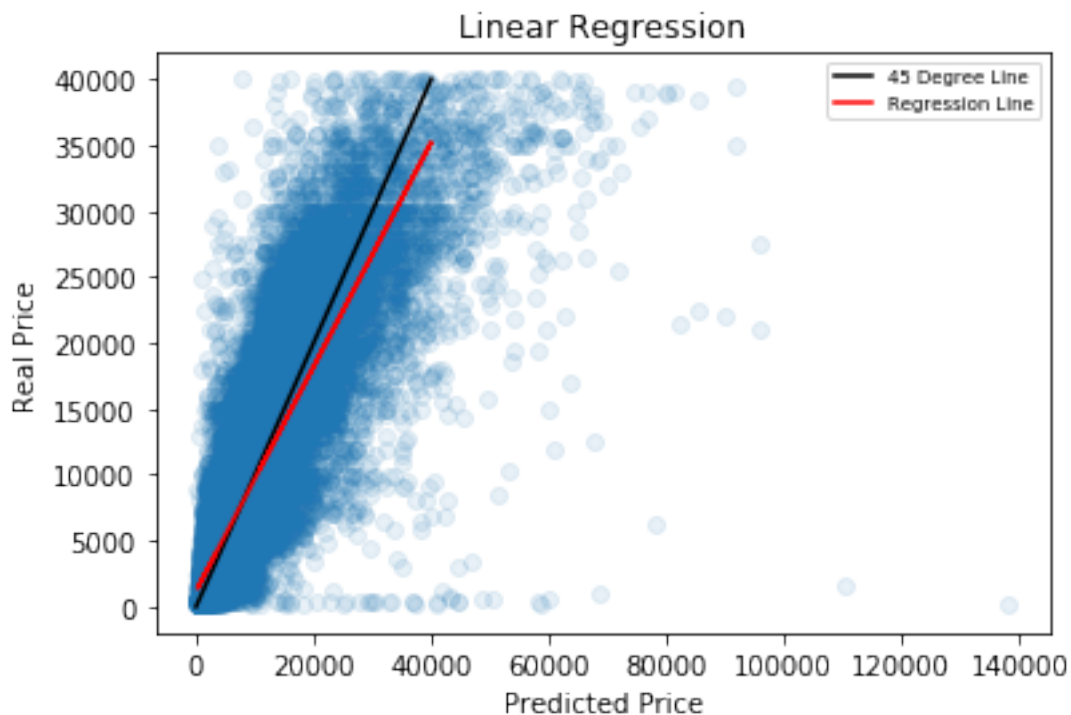
We used Linear Model, K Nearest Neighbor and Random Forest. The most successul one was Random Forest with 88.01% R-squared score. Followed by KNN with 83.70% and the worst one was as expected Linear Regression with 78.67%.

4.1.1 Linear Regression

Linear regression R-squared score was 78% but it overshoot for some cars and cause outliers also underpricing was very common. Despite 40k is maximum price in our dataset Linear Regression model made some predictions even 160k. Beside that the line was very thick so those predictions were not good.

We also tried to run Linear Regression without “model” variable to see the effects of it. It gave 3% less accuracy on the dataset. For Linear Regression full data gave better results.

```
[104]: plt.scatter(np.exp(y_predLR), np.exp(y_testLR.values), alpha=0.1)
plt.xlabel('Predicted Price')
plt.ylabel('Real Price')
m, b = np.polyfit(np.exp(y_predLR), np.exp(y_testLR.values), 1)
plt.plot([0, 40000], [0, 40000], color = 'black')
plt.plot(np.exp(y_testLR.values), m*np.exp(y_testLR.values) + b, 'red')
plt.legend(["45 Degree Line", "Regression Line"], fontsize = "x-small")
plt.title("Linear Regression")
plt.show()
```



```
[105]: print("Linear Regression R-squared score(with model):", scoreLR*100, "%")
print("Linear Regression R-squared score(without model):", scoreLRnm*100, "%")
```

Linear Regression R-squared score(with model): 78.58930520058853 %

Linear Regression R-squared score(without model): 75.69693562079023 %

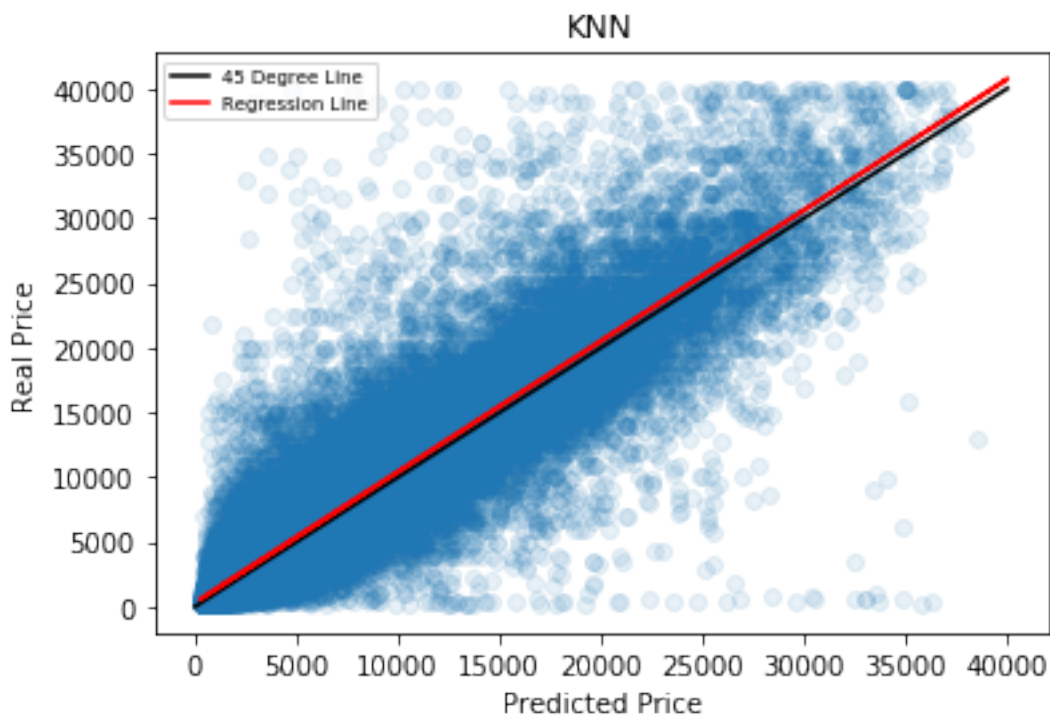
4.1.2 KNN

Since KNN is a Distance-Based algorithm we needed to start with scaling all variables. After data was ready next step was choosing how many neighbors to use. We decided to check one by one what are the mean squared errors and the R-squared score in the sample we took from data.

After running test on the sample we run the KNN with 6 neighbors for all data, results were relatively better. And as usual we fit the KNN with data without “model” and this time results were similar.

Since predictions were same with the data with and without “model”, below plot is with “model”.

```
[106]: plt.scatter(np.exp(pred), np.exp(y_testknn.values), alpha=0.1)
plt.xlabel('Predicted Price')
plt.ylabel('Real Price')
m, b = np.polyfit(np.exp(pred), np.exp(y_testknn.values), 1)
plt.plot([0, 40000], [0, 40000], color = 'black')
plt.plot(np.exp(y_testknn.values), m*np.exp(y_testknn.values) + b, 'red')
plt.legend(["45 Degree Line", "Regression Line"], fontsize = "x-small")
plt.title("KNN")
plt.show()
```



```
[107]: print("KNN R-squared score(with model):", accur*100, "%")
print("KNN R-squared score(without model):", accurNM*100, "%")
```

KNN R-squared score(with model): 83.6361413985903 %

KNN R-squared score(without model): 83.60531864475668 %

4.1.3 Random Forest

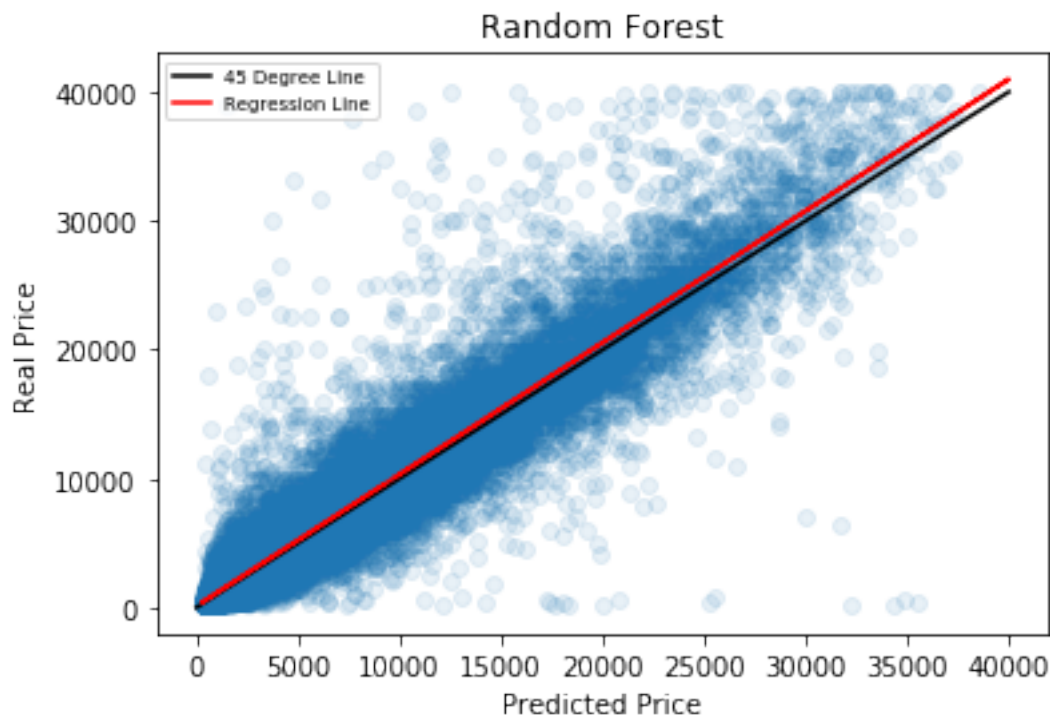
Last method we used was Random forest. We changed the way of encoding the data for this method in order to decrease depth of trees shorter.

We used KFold method to split our data to be able get more from it and see if R-squared score changes depending on the which part of data we are training and testing our model.

Results were better than Linear Regression and KNN.

And as usual we fit the KNN with data without “model” and this time R-squared score results were same. Although in without “model” prediction plot we realized that the line was thicker than predictions with “model”. So we can say that thanks to “model” in the data errors were smaller. For Random Forest data with “model” was better that’s why as a final plot belongs to that.

```
[108]: plt.scatter(np.exp(predRF) ,np.exp(y_testrf), alpha=0.1)
plt.xlabel('Predicted Price')
plt.ylabel('Real Price')
m, b = np.polyfit(np.exp(predRF), np.exp(y_testrf), 1)
plt.plot([0, 40000], [0, 40000], color = 'black')
plt.plot(np.exp(y_testrf.values), m*np.exp(y_testrf.values) + b, 'red')
plt.legend(["45 Degree Line","Regression Line"], fontsize = "x-small")
plt.title("Random Forest")
plt.show()
```



```
[109]: print("Random Forest R-squared score(with model):", scoreRF*100, "%")  
       print("Random Forest R-squared score(without model):", ttsscoreNM*100, "%")
```

Random Forest R-squared score(with model): 88.02434805259765 %

Random Forest R-squared score(without model): 87.26288219023658 %

Random Forest(with model) gave the best result 88%. For this dataset Random Forest perform in a best way in prediction and also computational power and time. We are thinking that it is the best fit among all 3 methods we used.

Despite the quality of the dataset we believe that results were pretty outstanding.